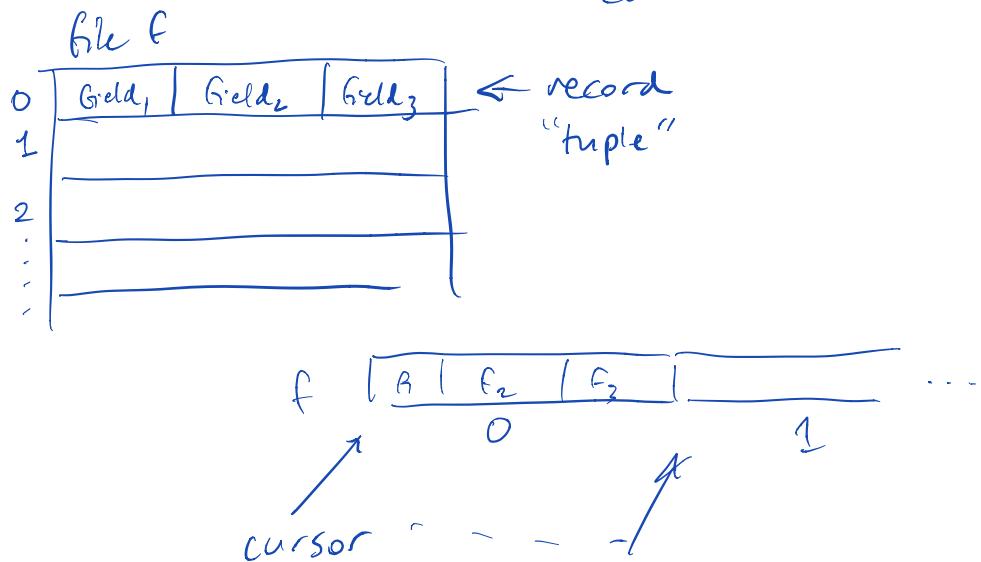


emery@cs.umass.edu

SQL - Structured Query Language

Container classes

```
set<T>    vector<T>
list<T>      vector<int> x;
:           x[4] = 12;
cout << x[4] << endl;
```



begin
next field
next record

Salary database Name DNI Salary

```

d=dbopen ("salary.db");
cursorBegin(d);
while (cursorPosition(d) != cursorEnd(d)) {
    s=readField(d,3);
    if (s < 10,000) {
}
}

```

1M

procedural
imperative PLs

(Java, C++, Python, PHP,
Pascal)

1M entries

sorted
by salary ?

```

d=dbopen ("salary.db");
cursorBegin(d);
while (cursorPosition(d) != cursorEnd(d)) {
    s=readField(d,3);
    if (s < 10,000) {
        if sorted(d)
}
    } else { break; }
}

```

100 records

inflexible
hard to optimize
breaks abstraction

low selectivity

city = 'Oxford' near
and
salary > 20.000 Euros
100%

Census
EU 300M
each record 1MB

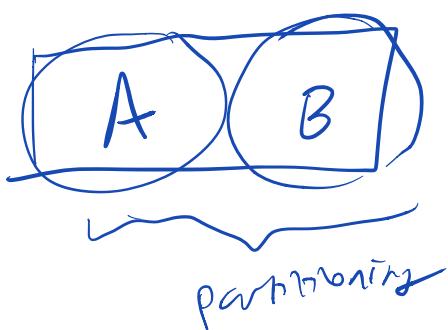
statistics
about DB

→ big perf benefit

in-memory 6000
but expensive

- only for medium workloads
(e.g. Not TB)
- "volatile"
 - not persistent
 - not durable

- indexing



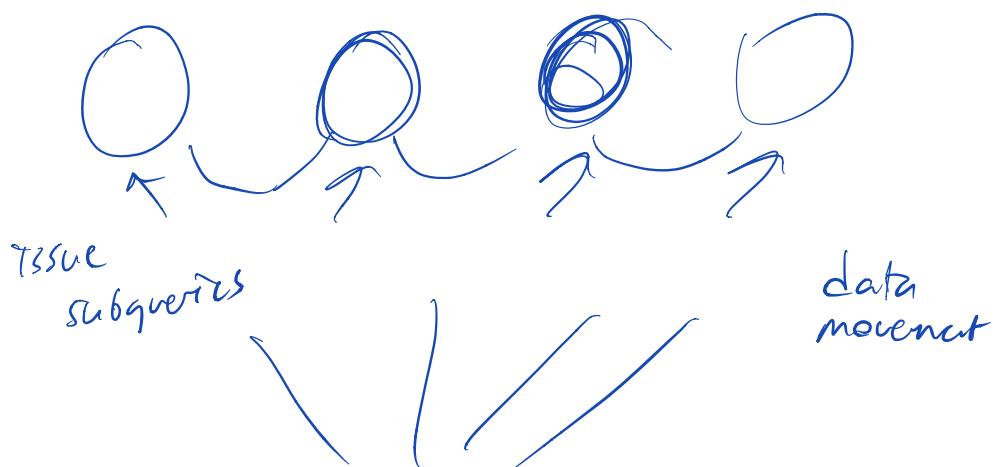
SQL — declarative

```
SELECT emp.name, emp.salary  
FROM emp  
WHERE emp.salary < 20000
```

query optimizer

syntax sugar
implicit parallelism

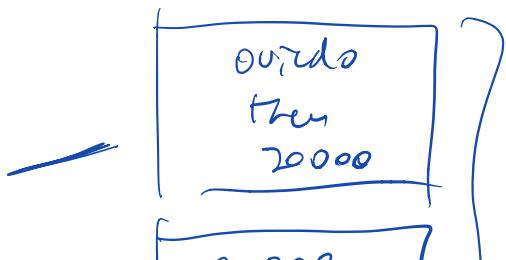
System R IBM (SQL) (DB2)



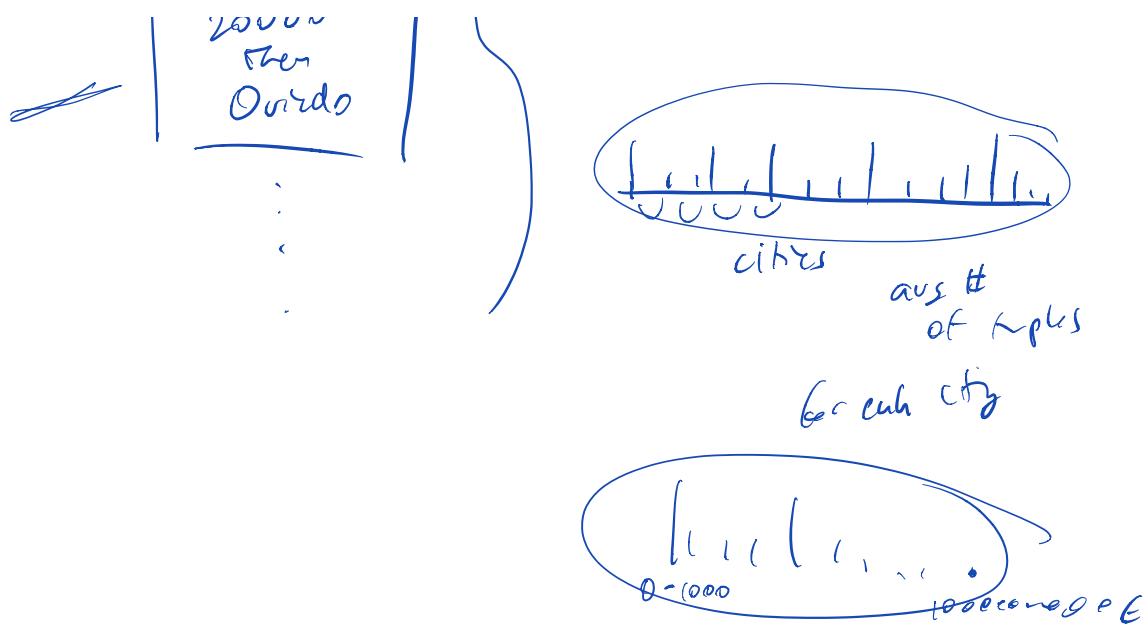
query plans

SELECT

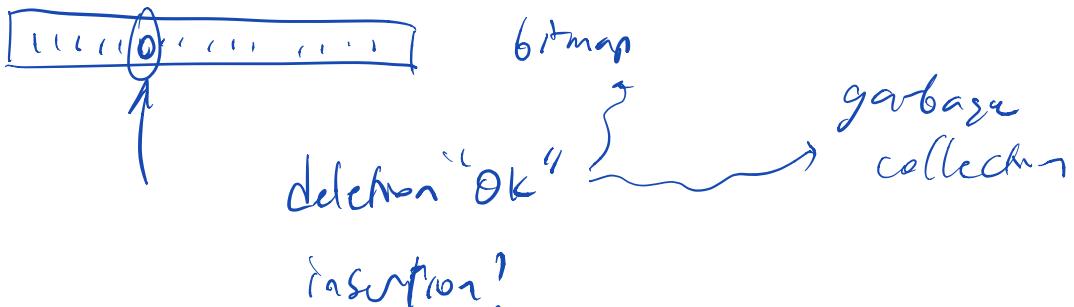
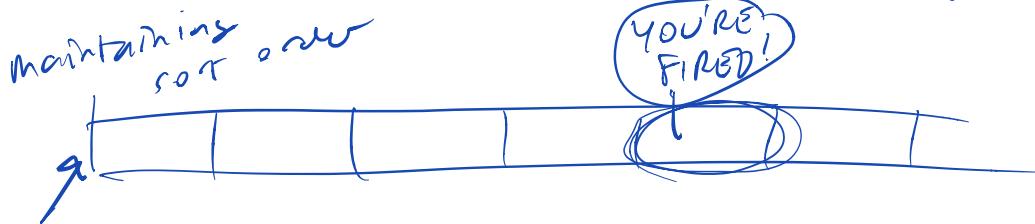
```
WHERE salary < 20000  
AND city = 'Orlando'
```



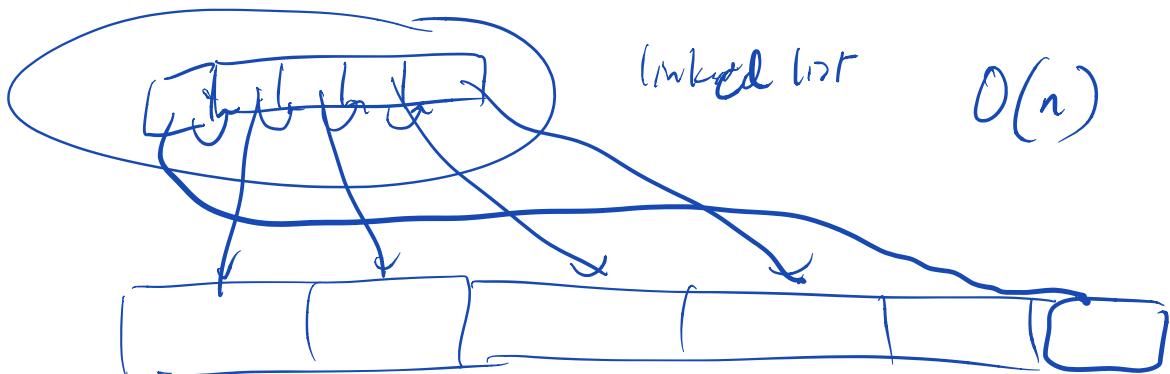
histogram by city



selectivity
cache queries
auto partitioning
leverage indexing



"all problems in CS
can be solved by adding
one level of indirection"

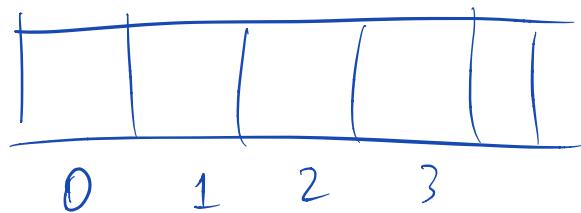


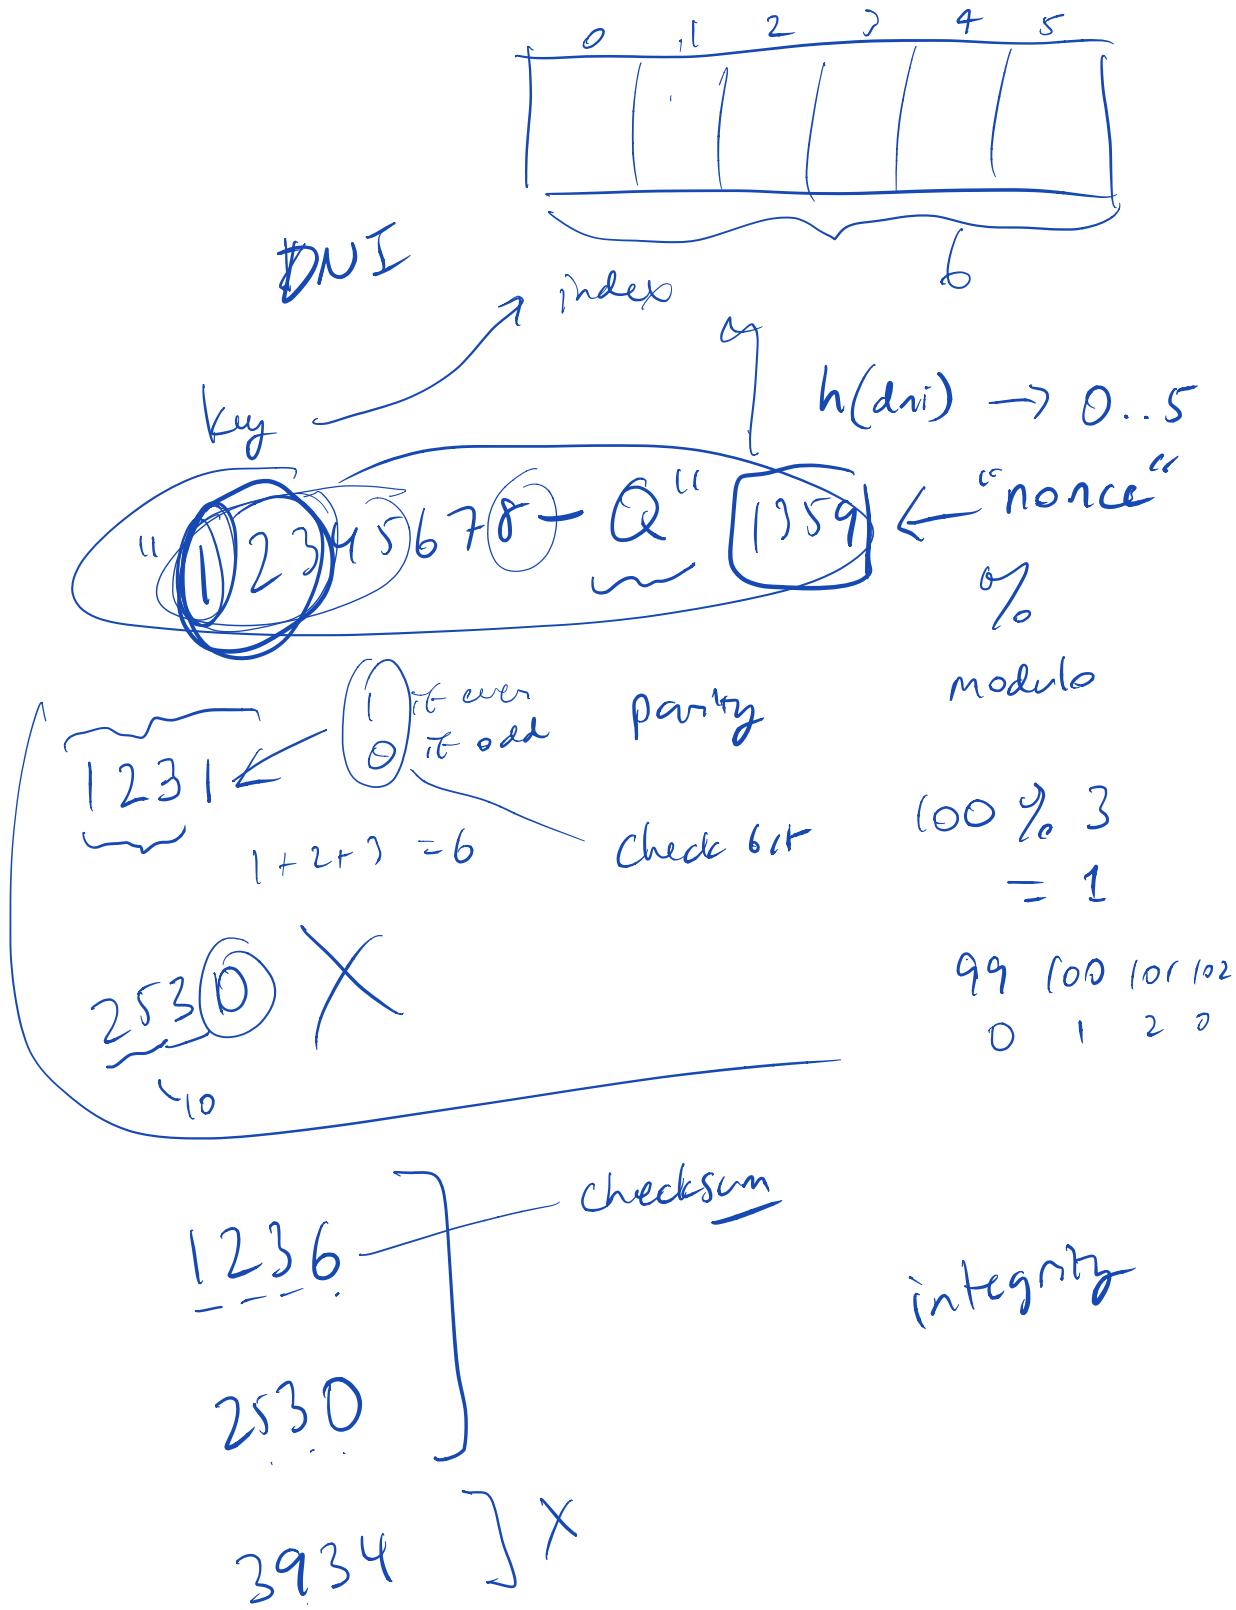
$O(\log n)$ — scalable

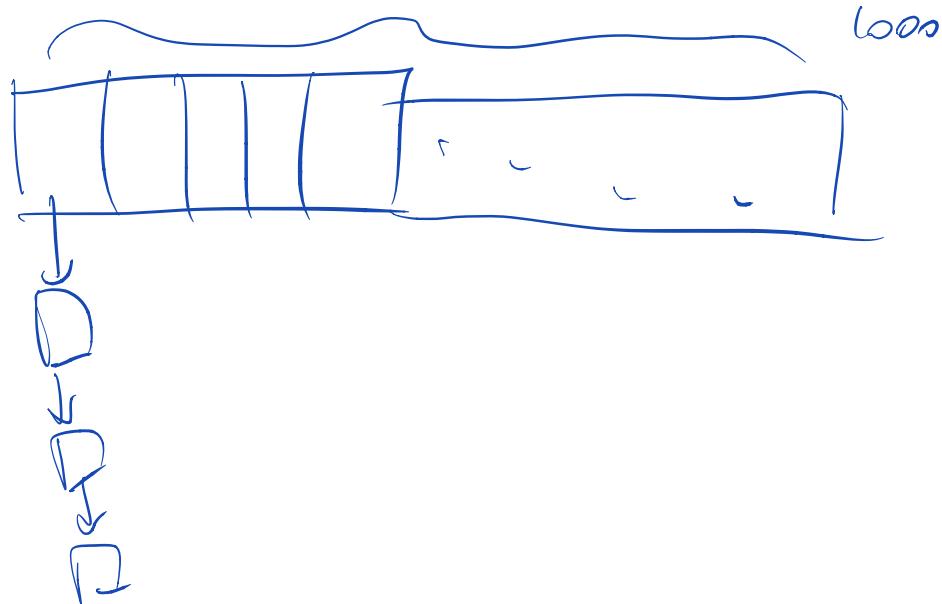
$O(1)$ — even better

hash table HashMap

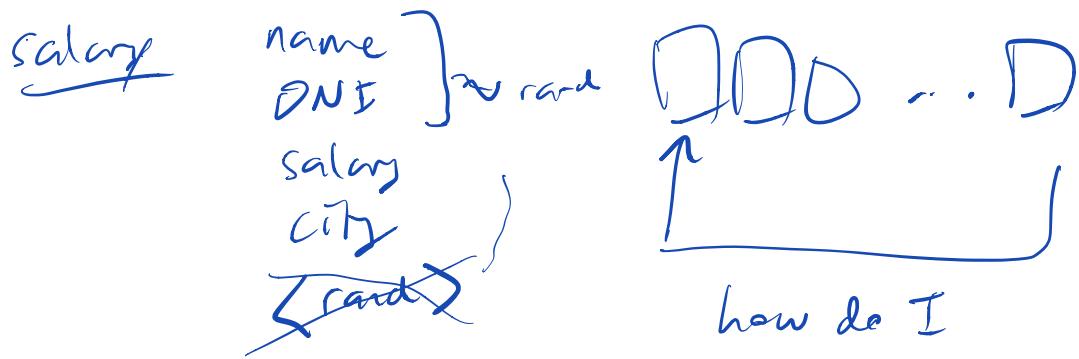
hash(key)







partitioning → goal: even distribution



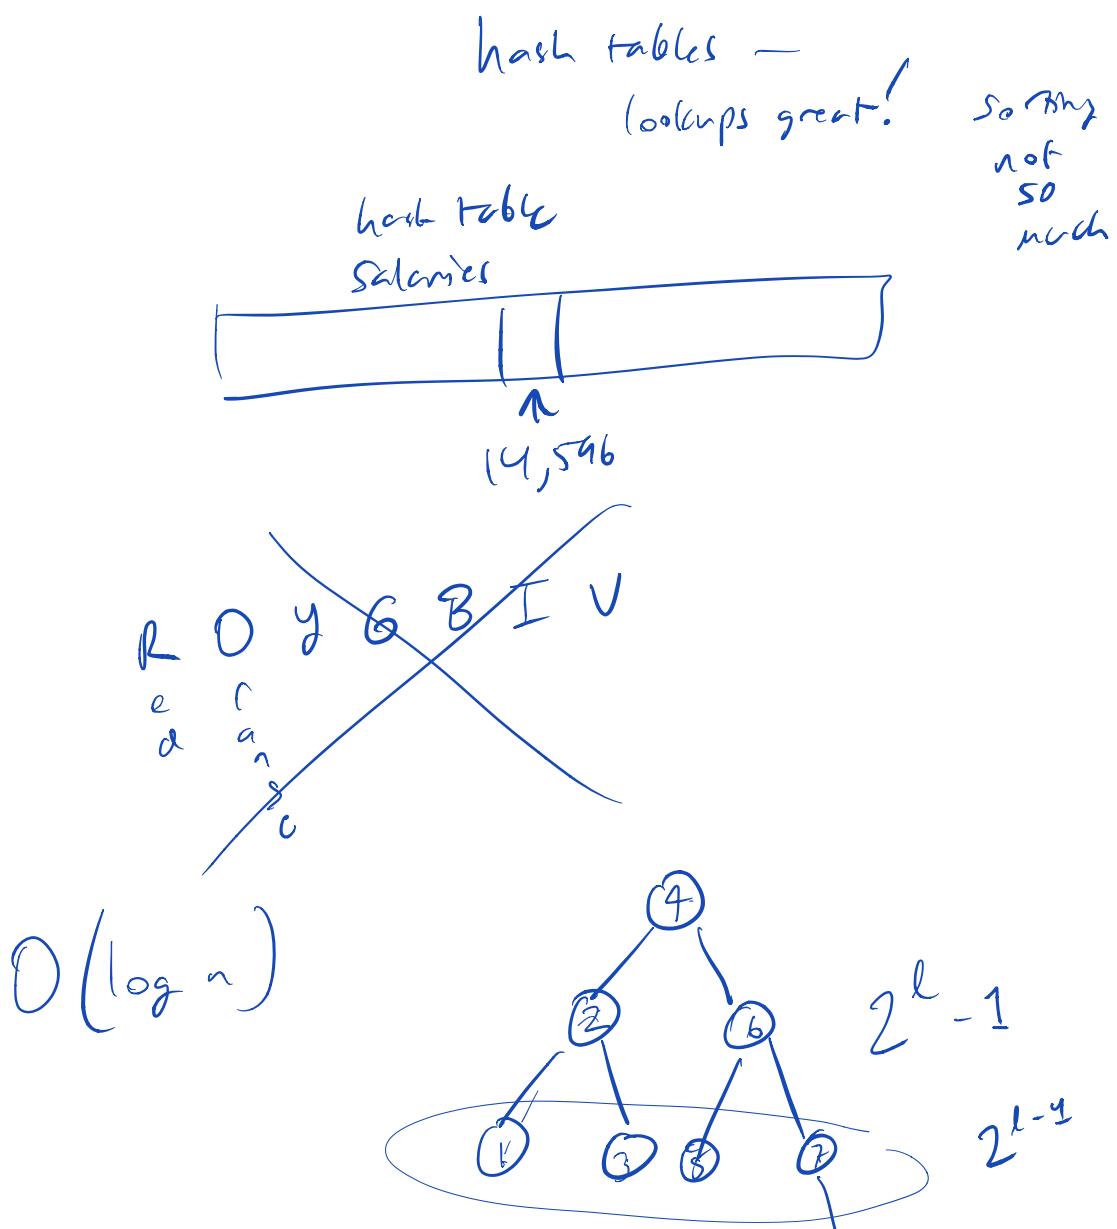
- access pattern?



hash partitioning

hashCode
trivx

Sorting linked lists = ~ sorted order
 no moving
 but $O(n)$

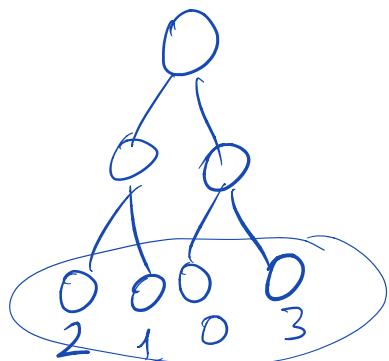


balanced (binary) tree

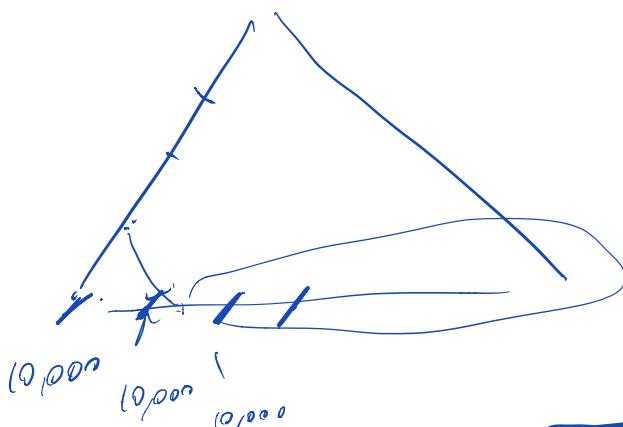


Red Black tree

AVL tree



10,001	9000	5	1,000,000
0	1	2	3

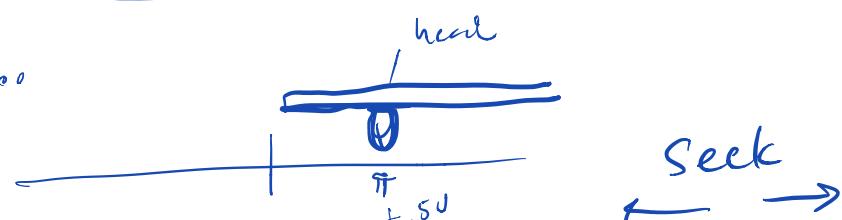


HDD

"Spinning media"
— conventional
disc

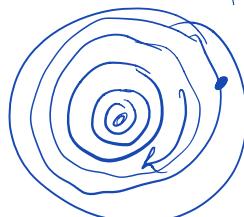
SSD

— Flash



seek latency

\gg
rotation latency ≈ 0



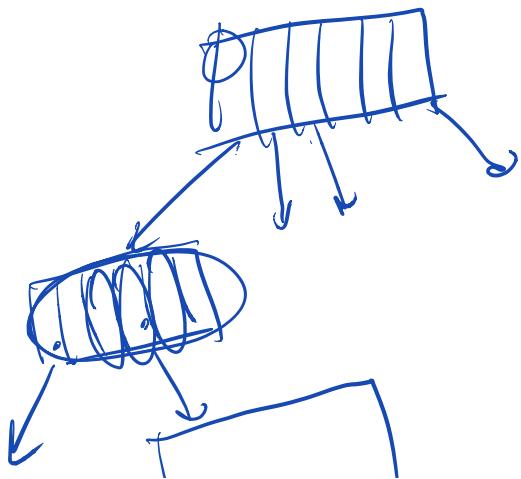
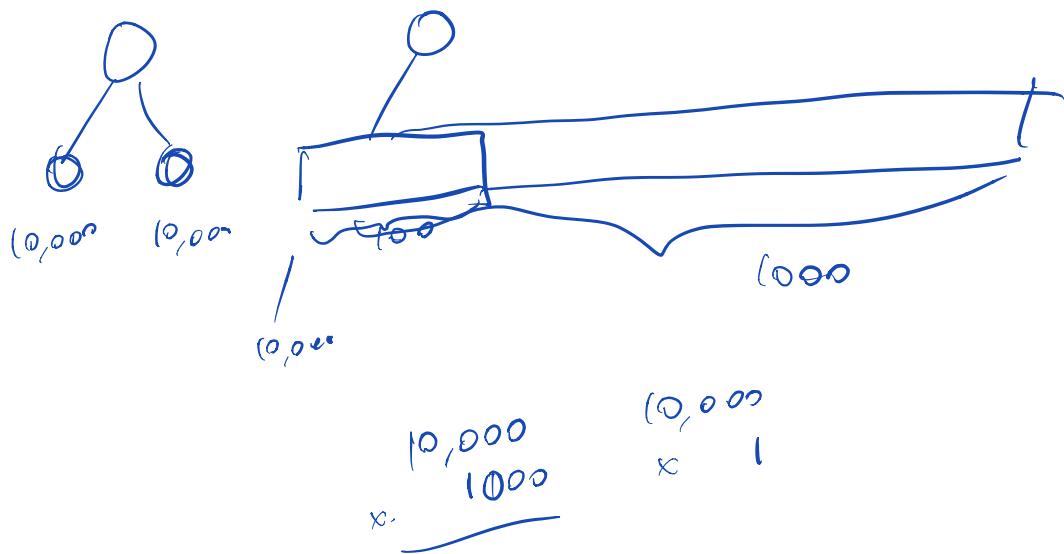
10,000 RPM

disk bandwidth
(height)

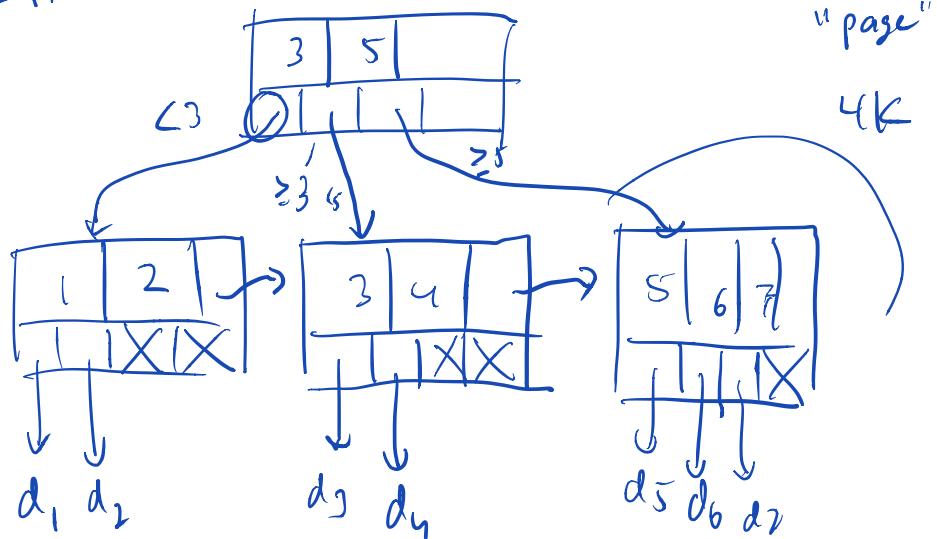
SSD: ~0 seek latency

$b_w < \text{disk } b_w$

cost $\sim 10\times$



B^+ -tree



index all fields?

- storage

- update (linearly w/ # of indices)

city, name

↑ ↑
primary key secondary key

SELECT

:
Group By city

"Sorted" By city

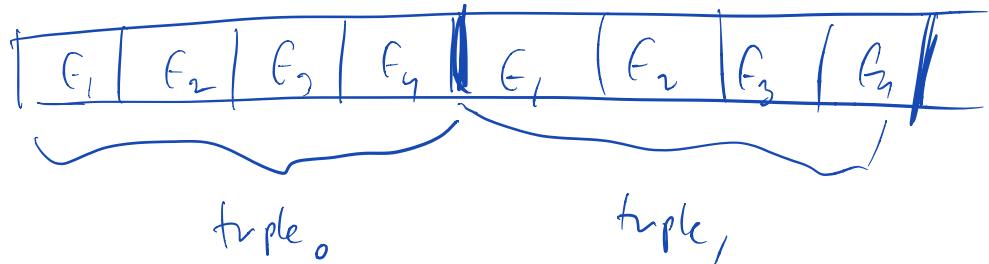
4 Fields ~ 4 indices

$$\text{pairs? } \binom{4}{2} = \frac{4!}{2!2!} = \frac{4 \cdot 3 \cdot 2!}{2 \cdot 1 \cdot 2!} \approx 6$$

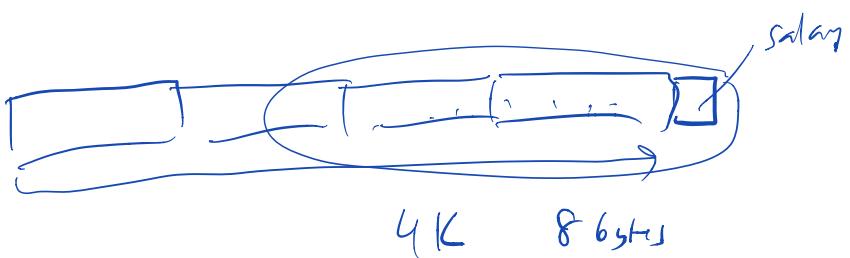
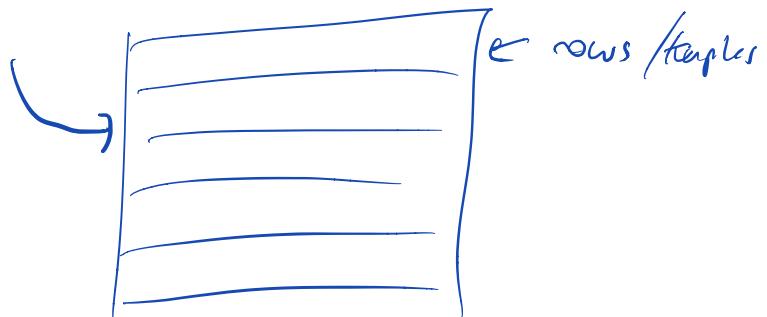
100 fields — 100 indices

$$\binom{100}{2} \quad \frac{100!}{2^{100}} \quad \text{or} \quad \frac{100 \cdot 99}{2} \approx 5,000$$

DB administrator



row-based layout

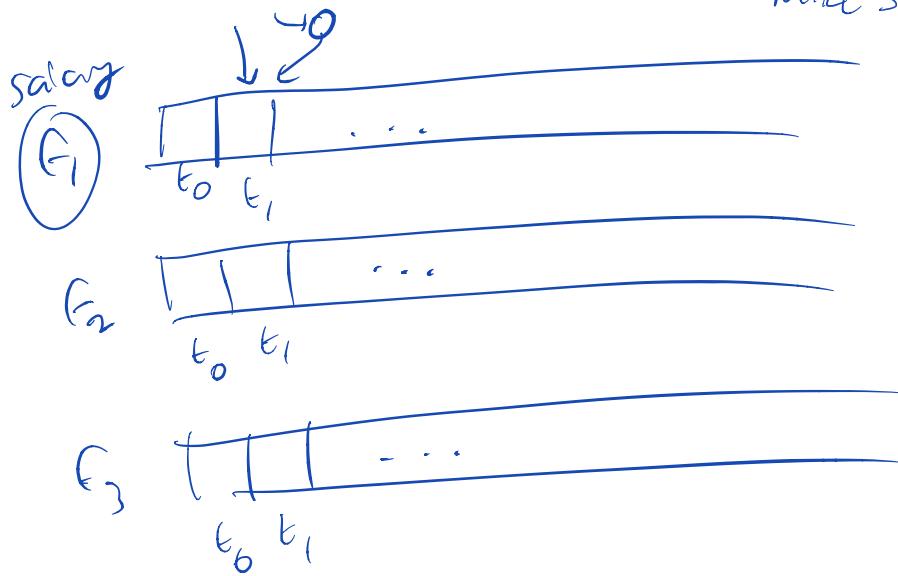


Column-based layout

Veniza

Mike Stonebraker

MIT



SIGMOD

VLDB