

PLANE FLIGHT# FUEL -
 AIRBUS IBERIA 235 — l

RING
PLANE FLIGHT# RING }
 # of planes in
 RING n
 ≤ 3

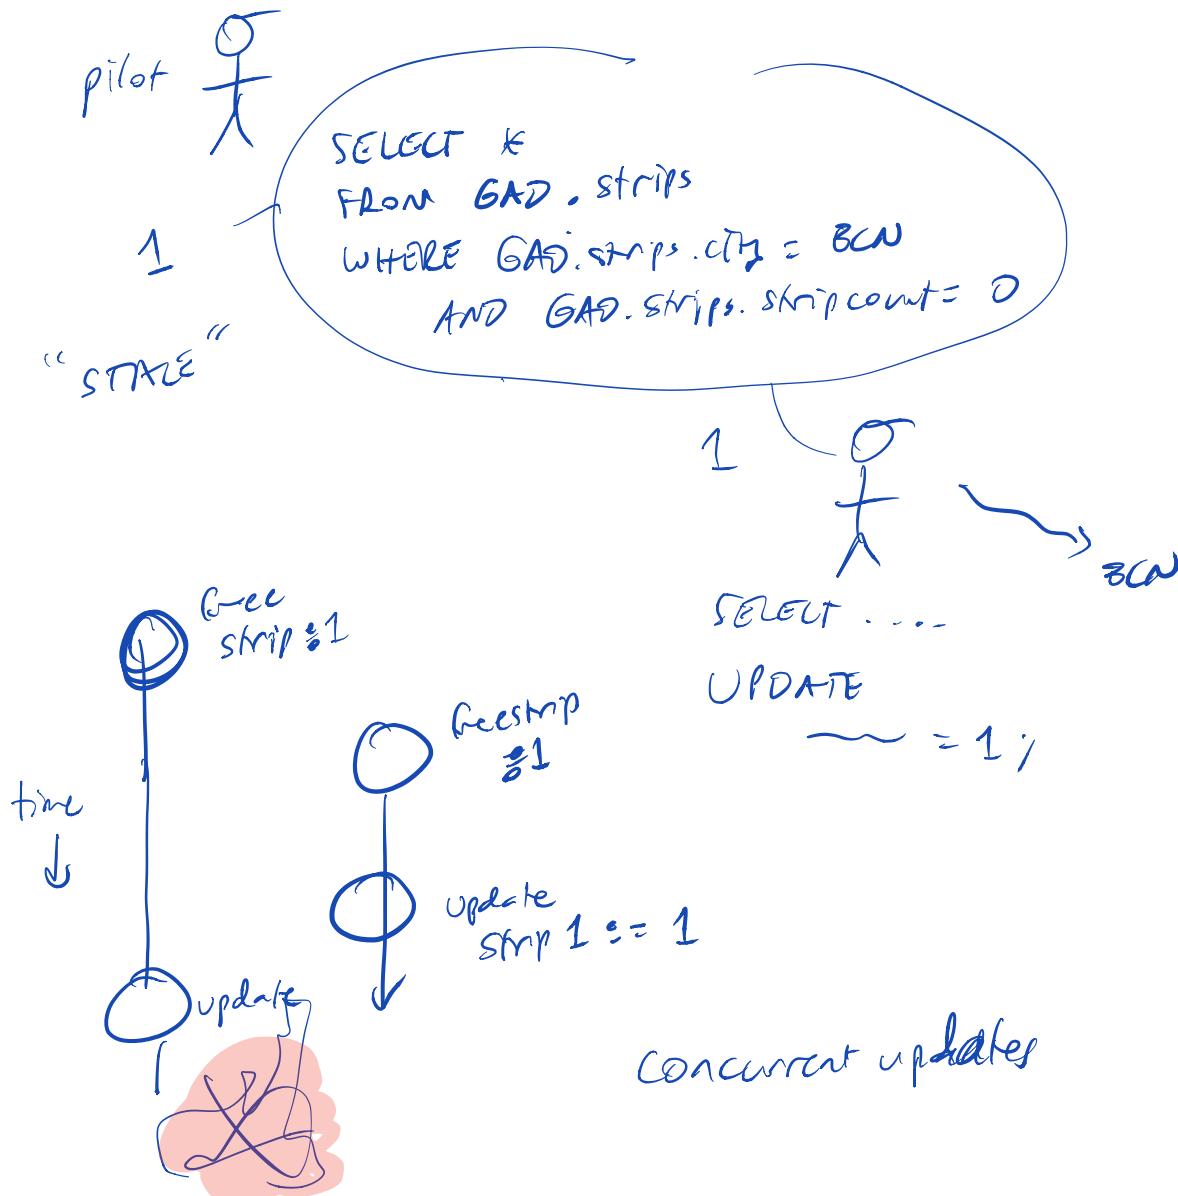
LANDING STRIP
 L a STRIP }
 # of planes
 STRIP n
 ≤ 1

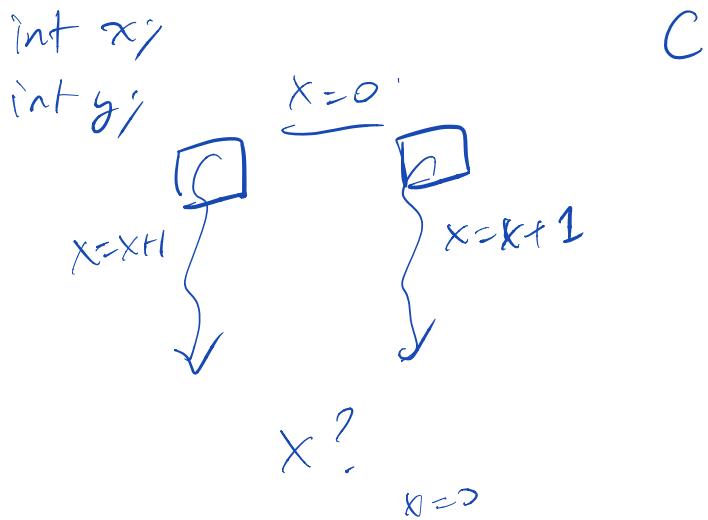
Global airport database

— concurrent updates

— "consistency guarantees"

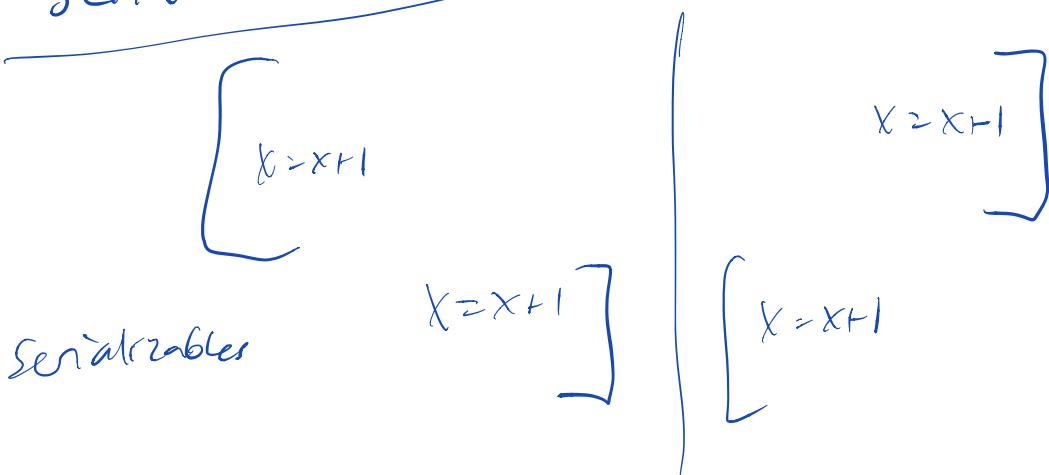
inconsistent =
broken invariants





$\text{LOAD } x, R\emptyset$ $\{ R\emptyset = x \}$ $\text{INC } R\emptyset$ $\{ R\emptyset = 1 \}$ $\text{STORE } R\emptyset, x$ $\{ x = 1 \}$	$\text{LOAD } x, R\emptyset$ $\{ R\emptyset' = x \}$ $\text{INC } R\emptyset$ $\{ R\emptyset' = 1 \}$ $\text{STORE } R\emptyset, x$ $\{ x = 1 \}$
--	--

Not
Serializable



SOLUTION (MT world)

→ LOCKS

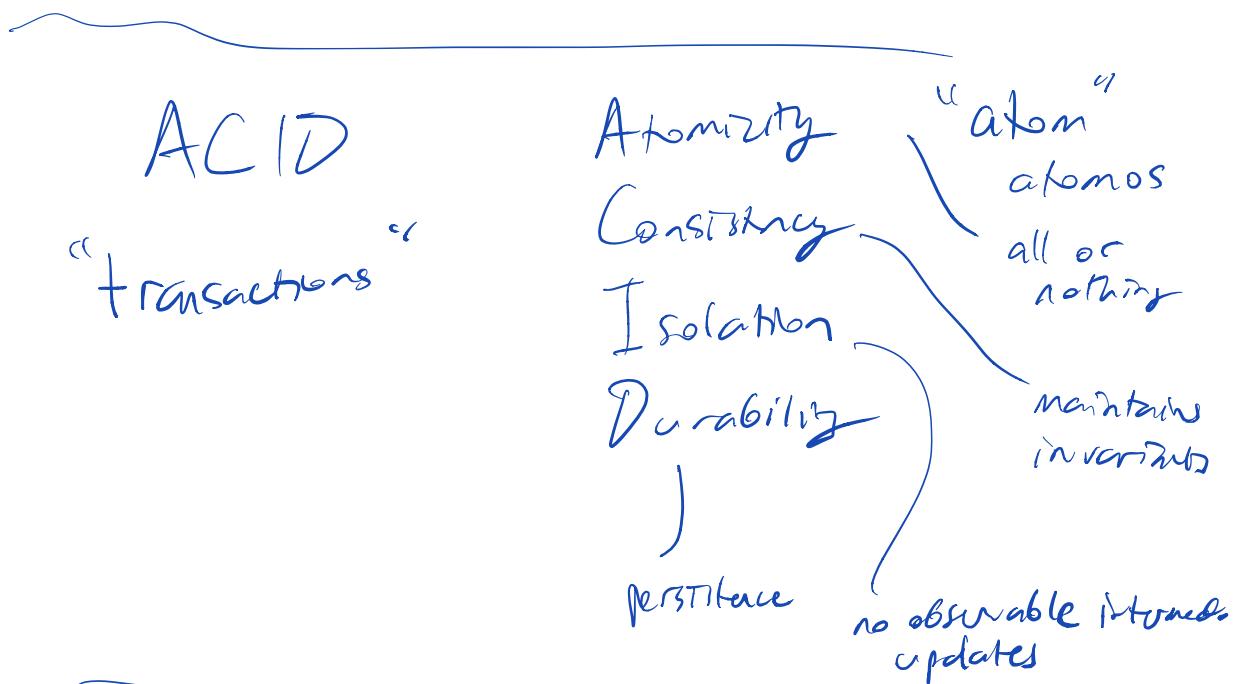
Mutual exclusion

Synchronized

`pthread_mutex_lock(& l);`

`x = x + 1;`

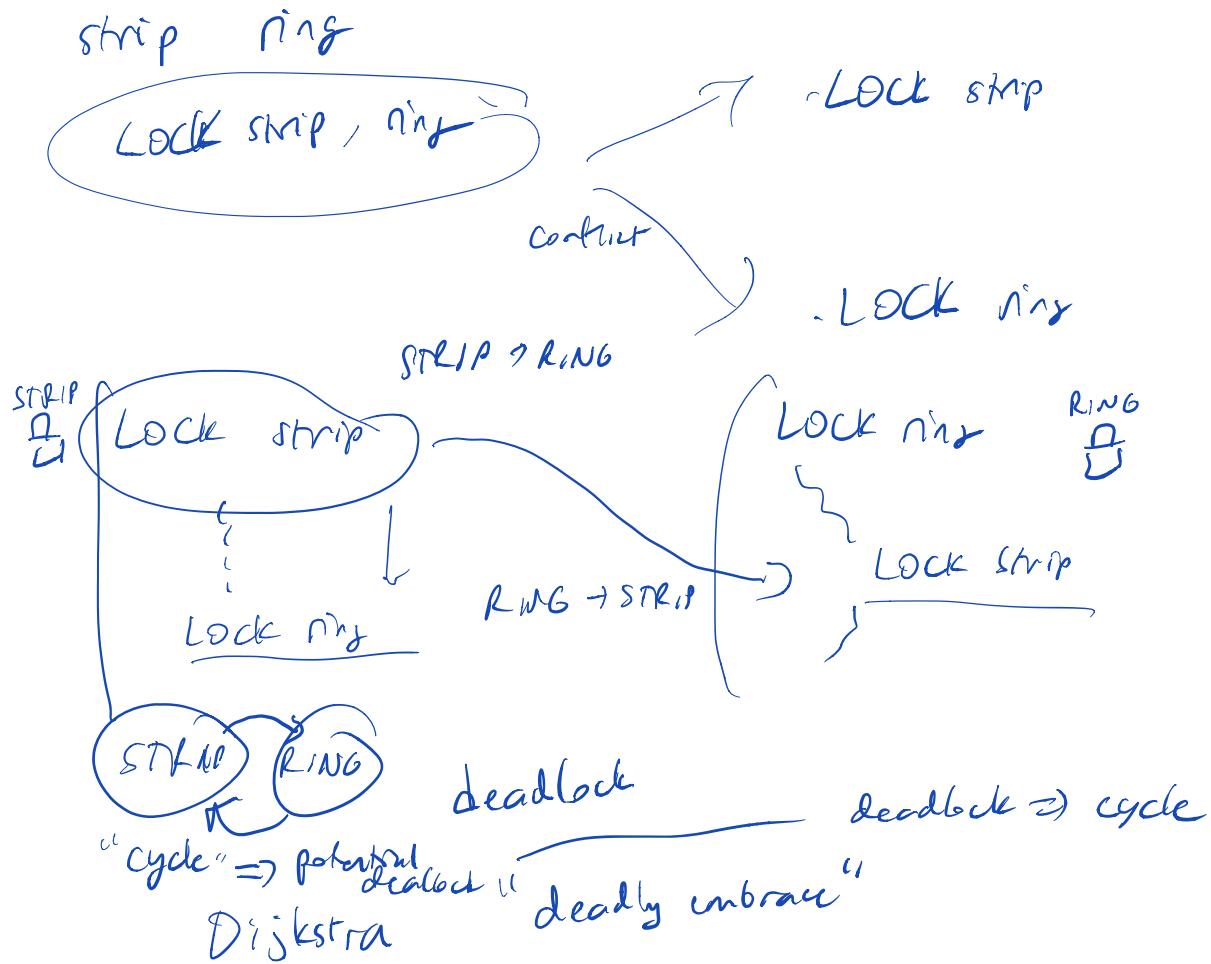
`pthread_mutex_unlock(& l);`



LOCK Strip
SELECT
UPDATE ...

UNLOCK Strip

→ "exclusive lock"



deadlock avoidance

request all locks
at start

deadlock detection

- ✓ { try all locks
 - { If one is locked
 - retry all \rightsquigarrow spinning
- ✓ { "canonical" order
 - $R, S \quad S, R \Rightarrow "R, S"$ acquire

correctness ++! (0/0)

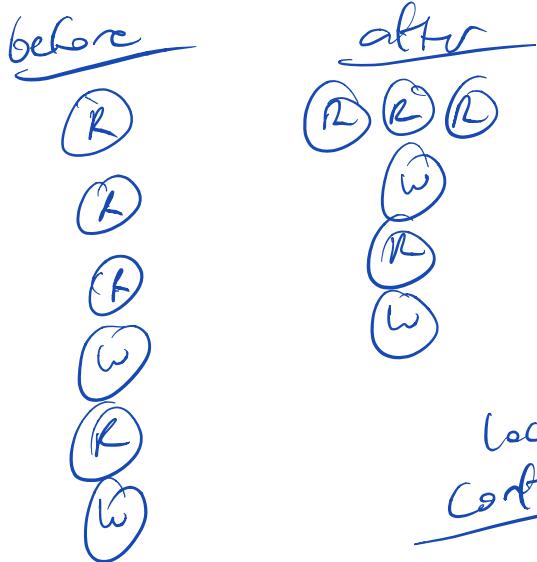
performance ?? --

"S,R" release
(reverse)

"database-level locks"

"READER" lock
Shared Lock } all readers
can operate concurrently

"WRITER" lock
Exclusive lock } only one writer (& no readers)
at a time.



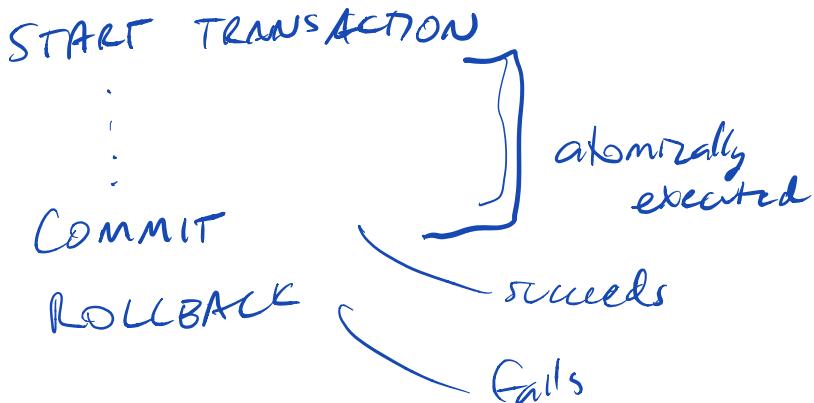
lock conflict table

function

	READ	WRITE
READ SHARED	✓	X
WRITE EXCLUSIVE	X	X

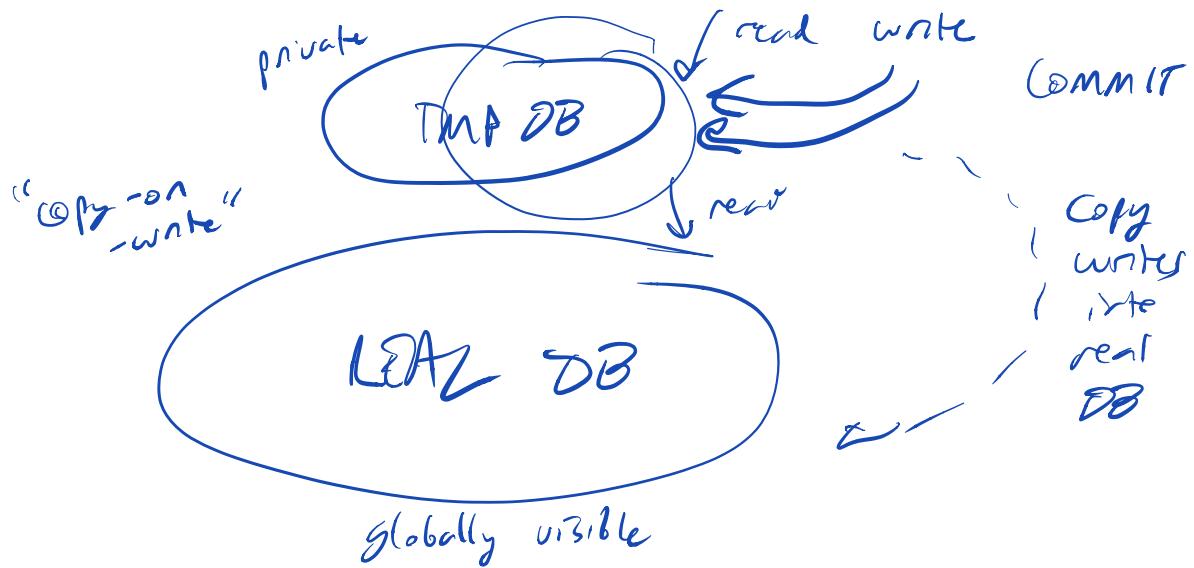
	DOL	DBW	RR	RW
DB READ	✓	X	(1)	
DB WRITE	X	X		
ROW READ	(1)	X		
ROW WRITE	X	X		

"ring + skip ≤ 4"



KimK $\leftarrow t=1$
 Khlock $\leftarrow 1$
 Kylock $\leftarrow 1$
 Khrk $\leftarrow t=1$
 If ($--Kardashians > 1$)
 ROLLBACK.

log
 KimK 0
 khlock 0
 kylock 0
 khrk 1
 "vado"



LOCKS : X ACTIONS

- Serializability ACD
- reasonable concurrency

(lots of) serialized xactions (esp. if long)

availability

"nines" of avail

99%

3.65 days unavailable

99.9% ~ 8 hours/year

99.999% ~ 5 min.

not just locking
failures

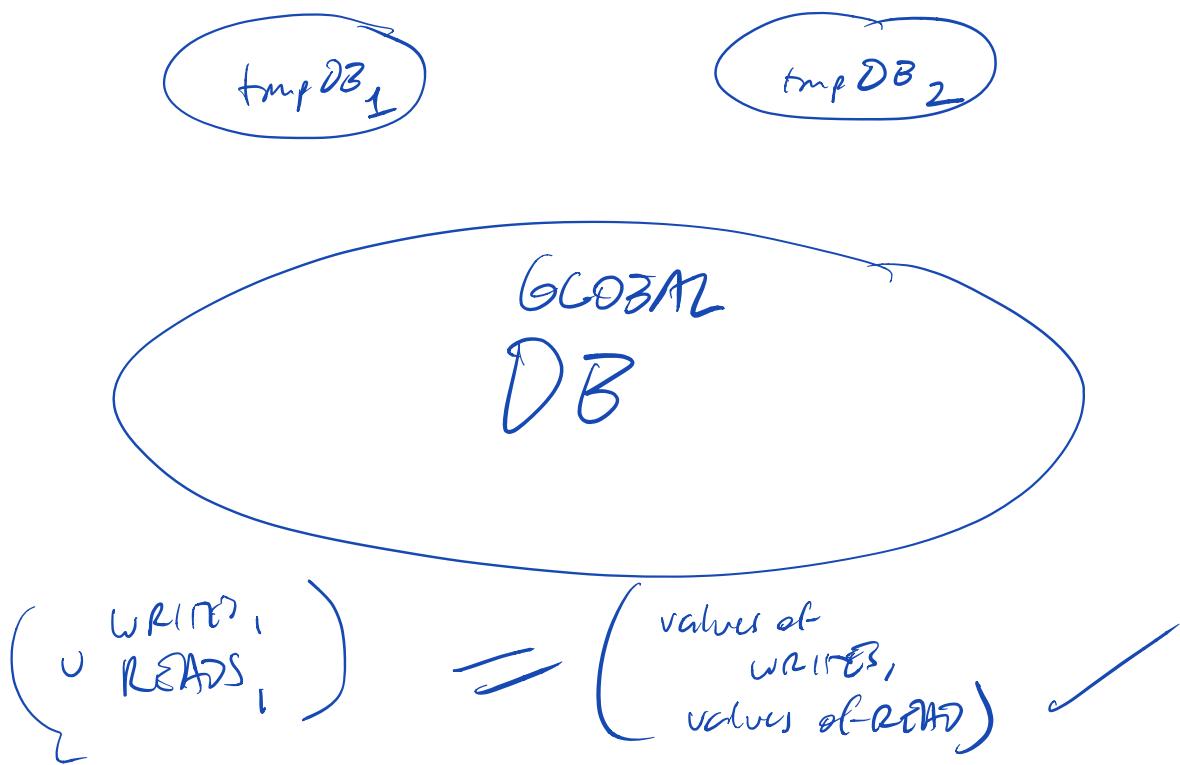
OLTP

online
transaction processing

mostly
write

BANK

"inconsistency" (Amazon)



Check read set & write set
 make sure not modified
 if not → commit
 if modified → ROLLBACK

Optimistic Concurrency

control
— "livelock"

Kung
& Robertson

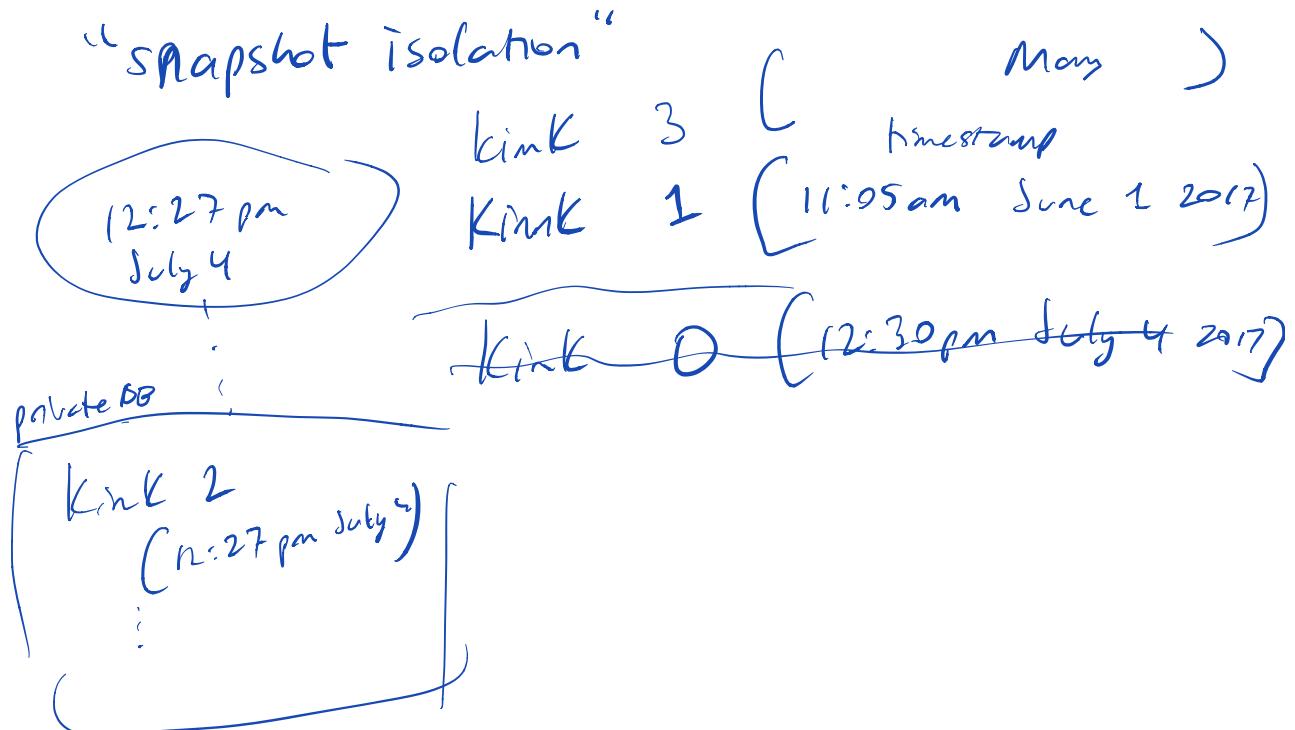
locking — pessimism
— progress

“unfair”

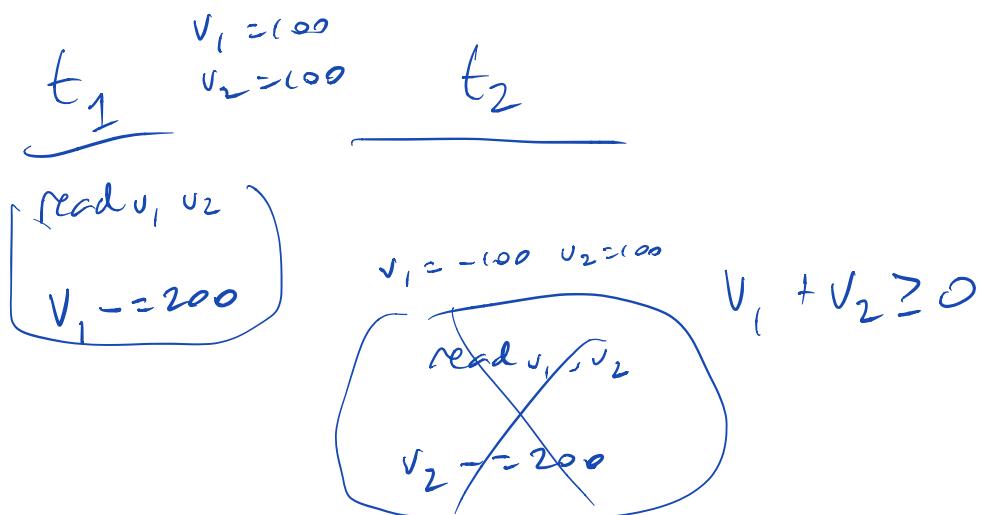
guarantees
serializability

aborts

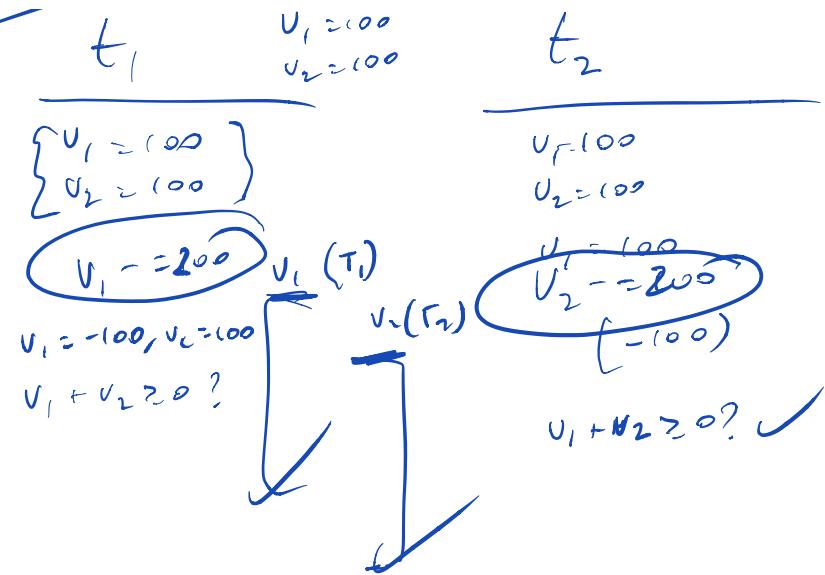
Multi-version Concurrency Control



write skew anomaly



Snapshots



$$U_1 = -100 \quad U_2 = -100$$

MVCC (snapshot isolation \Rightarrow more commits

some weirdness
 (write skew)
 "anomaly"

port.
 exclusive
 lock
 table

→ readers
 shared
 reads

→ row
 lock \rightarrow OCL \rightarrow MUCC

~~lock~~
 fault tolerance
 \Rightarrow replication

"serializing"

Sof

single
 point
 of
 failure



MTTF = Mean time to failure
availability

$$P[\text{fail.}] = p \\ /_{1000}$$

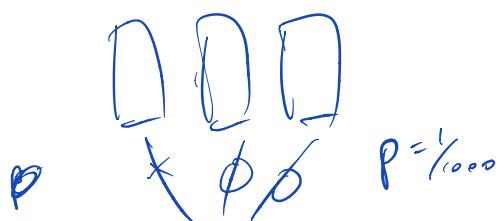


$$E(\# \text{ of com App.}) = \frac{1}{p} = 1000$$

$$\frac{1}{p=1/1000} \quad \frac{2}{p^2 \downarrow 1/100,000} \quad \frac{3}{p^3 \downarrow 1/10^9} \quad \dots$$

trimodular redundancy

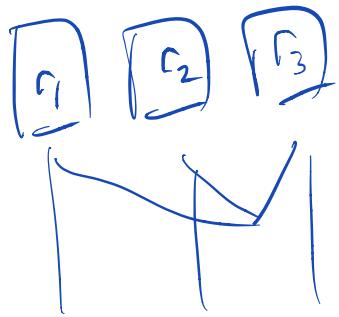
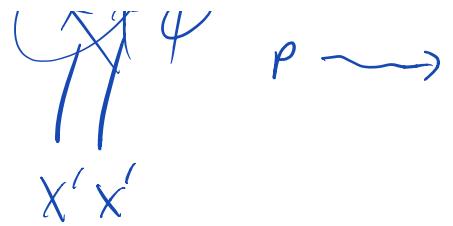
fail stop



Lamport
Byz Generals
problem

BFT





Tandem

Jim Gray