

<https://github.com/emenybergu/>

Systems - DS - UPC - 2017

fail-stop vs. incorrect results

what about delays?

FLP result

Fischer, Lynch, Paterson

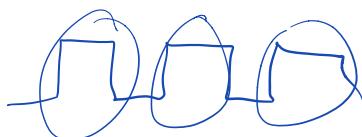
asynchronous systems vs. synchronous

distributed consensus
impossible

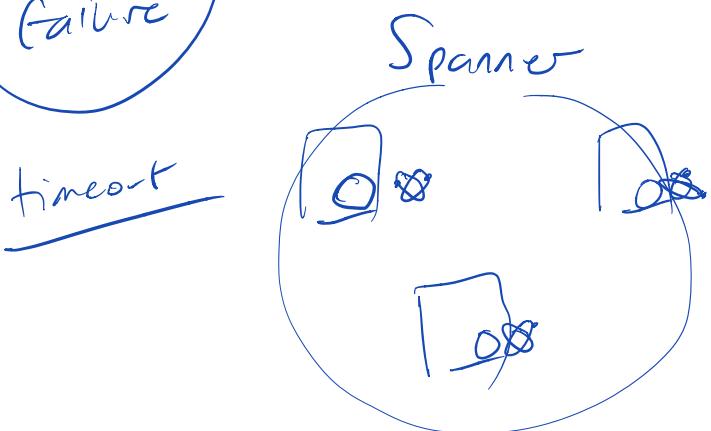
Clock skew/
NTP drift



□ □ □) "lock step"



requires a
clock



databases

- semantics serializability transactions

- query language

(DB admin - indexes)

— scalable? \Rightarrow (MapReduce vs. DBMS)

scale up — takes adv. of all cores on a node

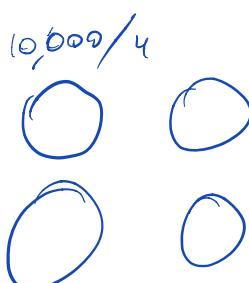
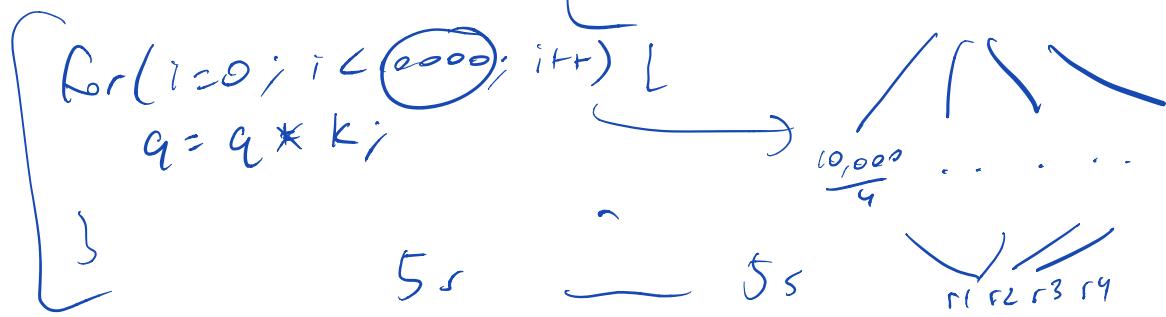
scale out — takes adv. of all nodes

P processes
 N nodes

$$\text{perf}(P \cdot 2) \approx 2 \cdot \text{perf}(P)$$

$$\text{perf}(N \cdot 2) \approx 2 \cdot \text{perf}(N)$$

linear scalability
no scalability
super linear scalability (!?)



automatic parallelization

→ explicit parallelism

(OpenMP)

→ MT

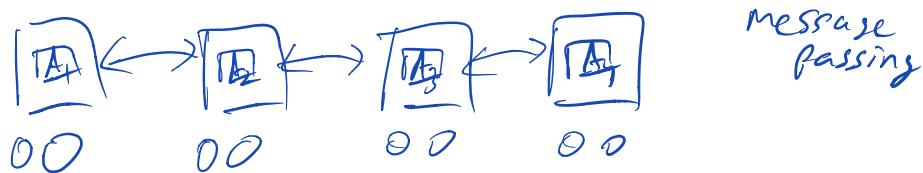
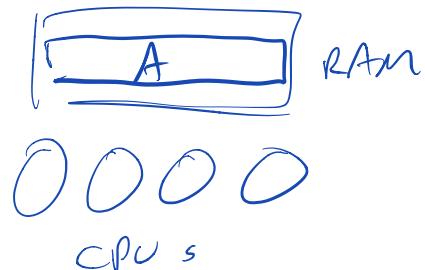
"embarrassingly parallel"

```

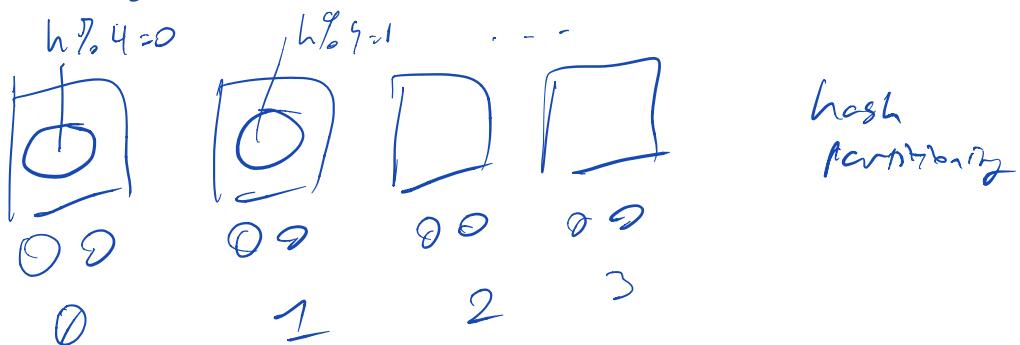
for( ) {
    x[i] ... x[i-1]
}

```

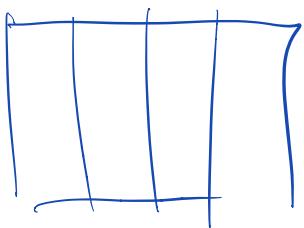
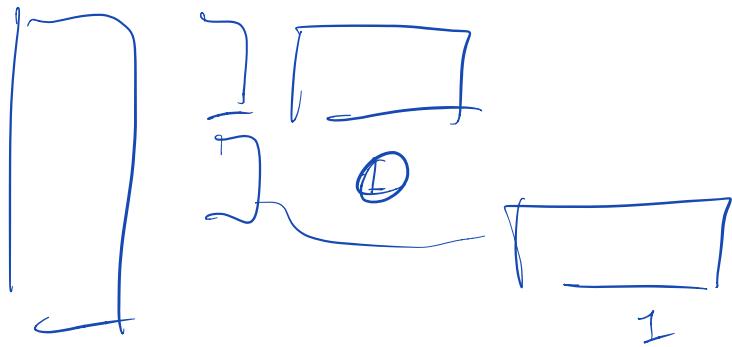
"shared memory system" multiple threads



database



Sharding



Scalability?

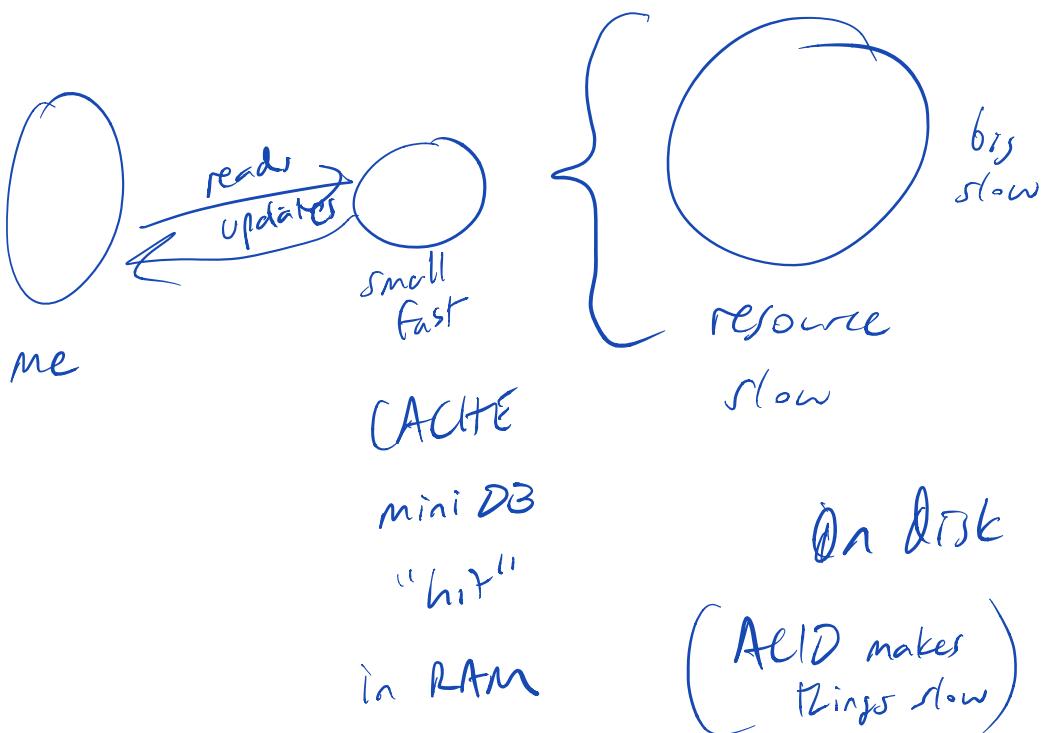
DB scale ~ 100

speed?

latency \rightarrow high

scalability

abstraction penalty



Cache API:

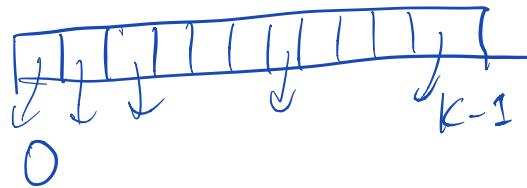
$v = \text{get}(\text{key})$	string	
$\text{set}(\text{key}, \text{value})$	dict	hash table
\ /	$v = x[\text{key}]$	
string	$x[\text{key}] = \text{value}$	

SHTA-1 (str) \rightarrow hash $\text{string} * \text{arr}[K];$

```

set(key) {
    h = hash(key)
    ind = h % K;
    return arr[ind];
}
set(key, value) {
    h = hash(key);
    ind = h % K;
    arr[ind] = value;
}

```



$\text{set}("Claudia", 10)$
 $\text{set}("Gordi", 10)$
 $\text{set}("Paula", 10)$



pigeonhole principle

\rightarrow collisions

get(key)

h = hash(key)

Ind: h > K

while (p != null)

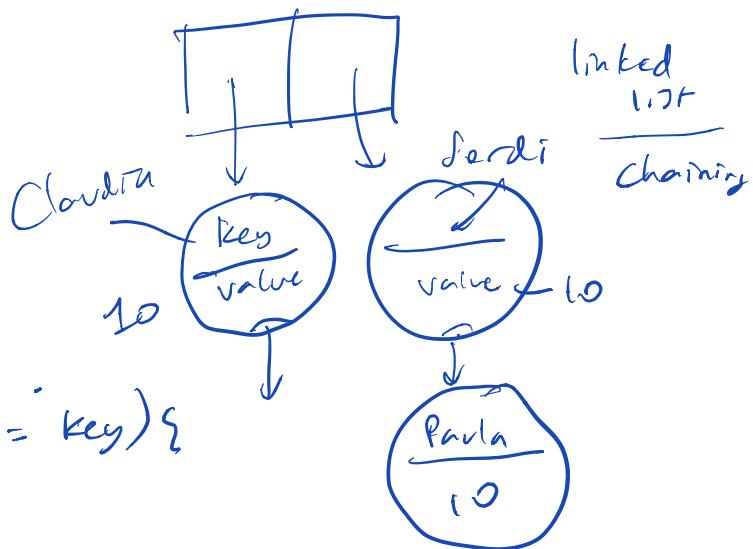
if (p->key == key) {

 p = p->next;

}

return p->value;

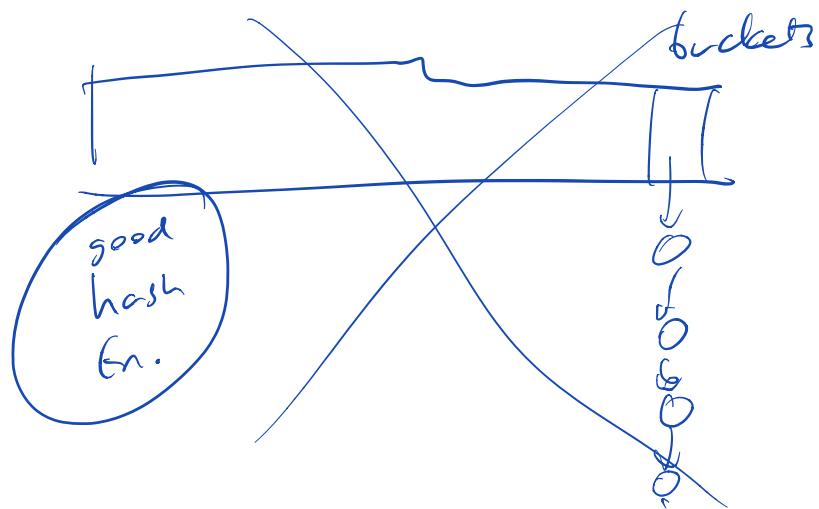
}



linked
list
Chaining

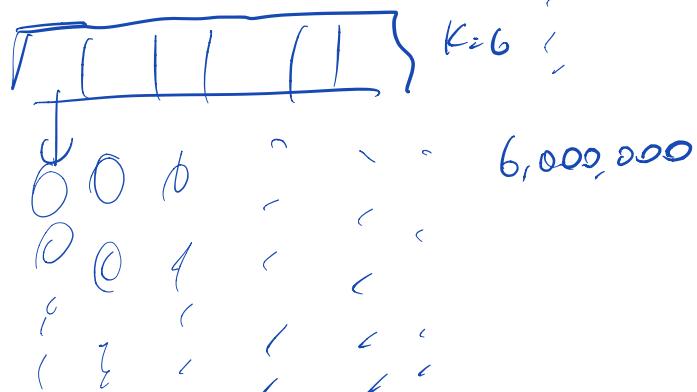
DoS
attack

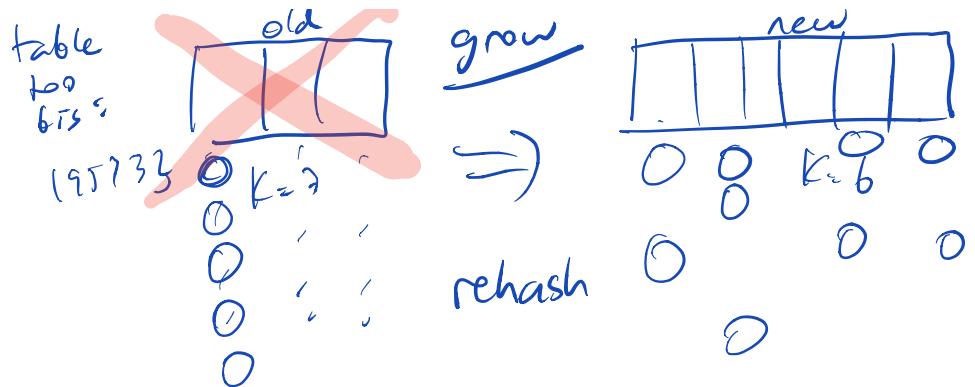
denial of
service



DDoS

distributed

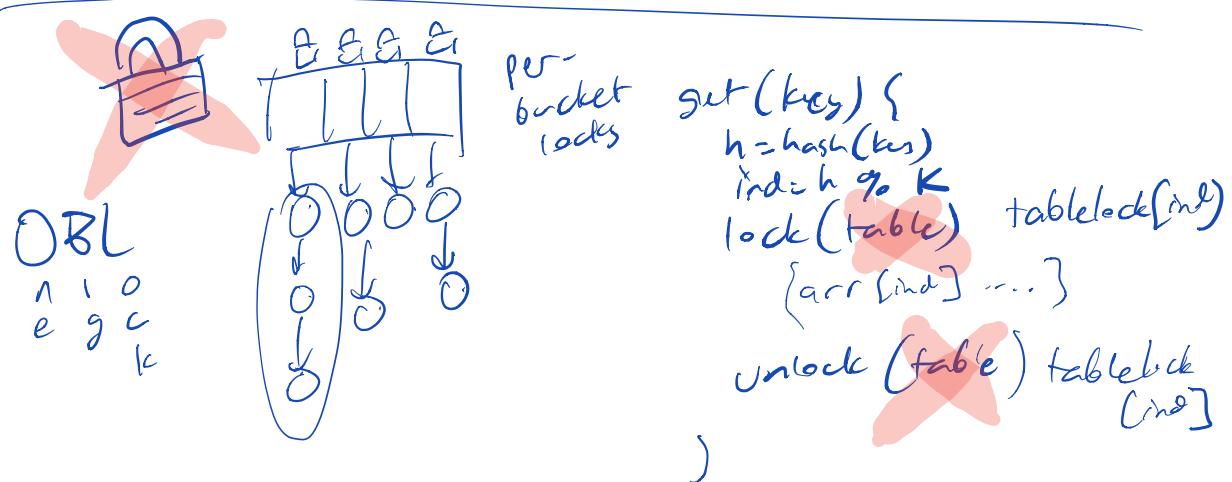




dynamiz hash table

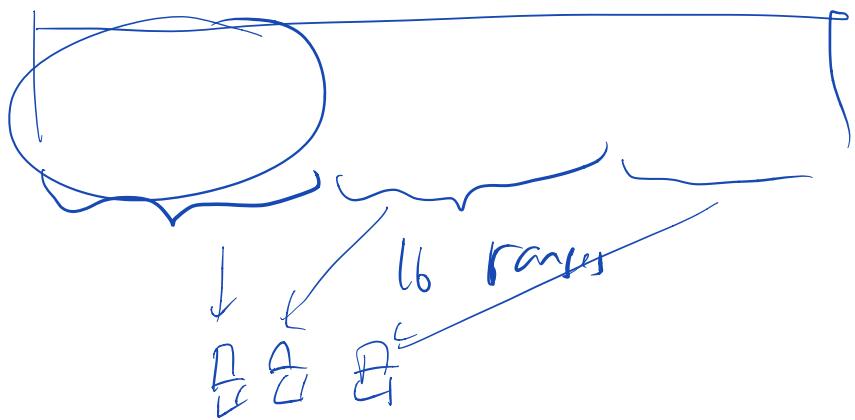
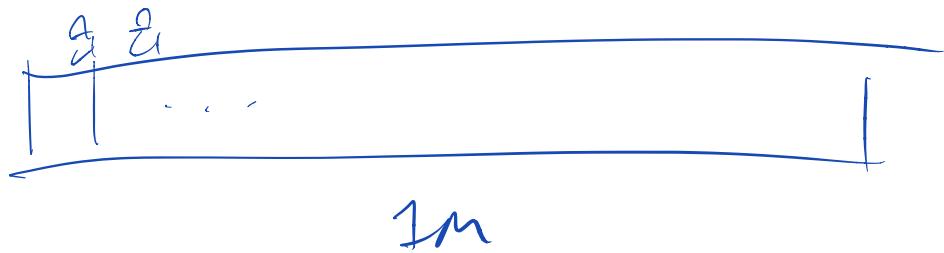
key /upc/systems/paula/ms1.ipd

value = 1 MB

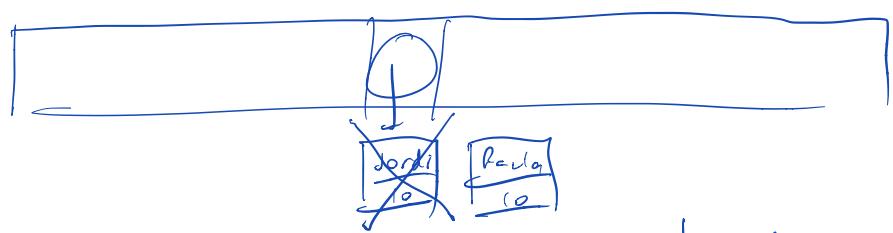
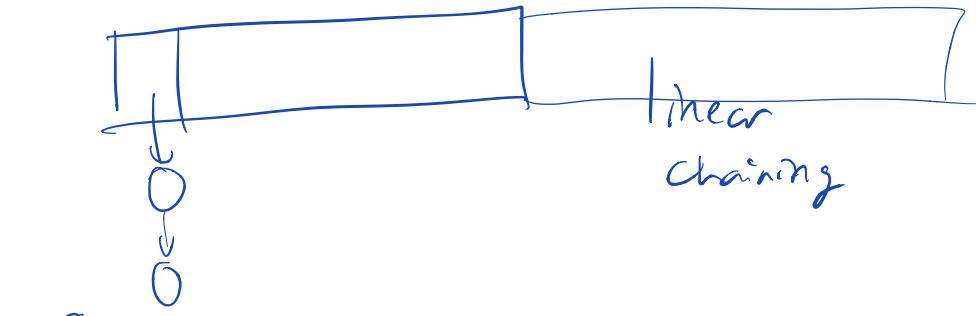


MTF
 move-to-front

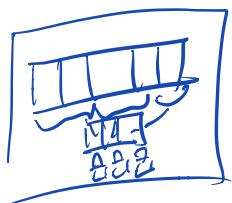
last item = LRU



Nicola's proposal

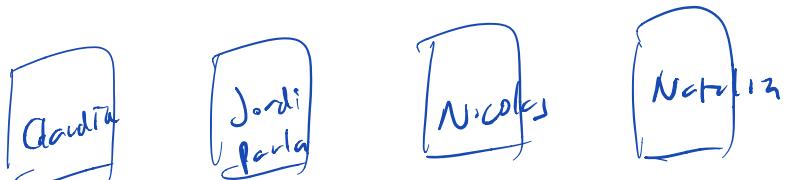


direct
mapped
(probably
not a win)



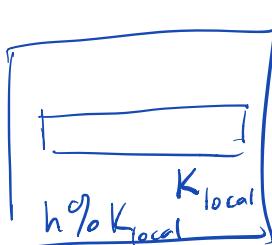
00000
CPUs

distributed hash table

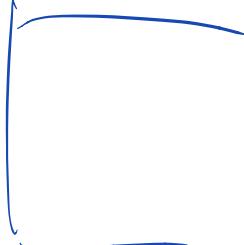
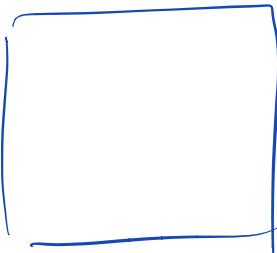


```
distrbget(key) {  
    h = hash(key)  
    machine = h % 4  
    v = machine.get(key)  
    return v;  
}
```

memcached



~

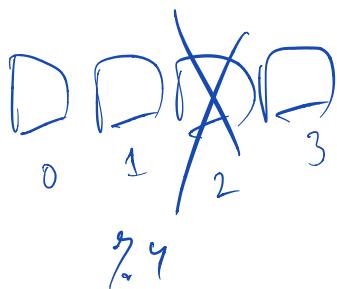
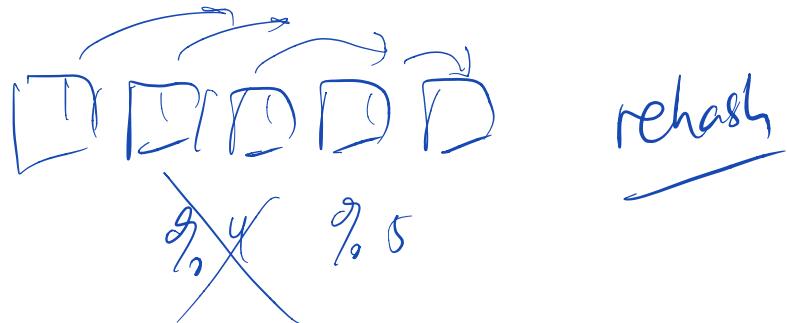


BIG

↑
h% K_{machines}

PROBLEM

adding or removing machines



$$P(\text{Failure}) = p \rightarrow .001$$

$$1 - \underbrace{(1-p)}_{\text{no failure}} \circ \underbrace{(1-p)}_{\text{no failure}}$$

$$1 - (1-p)^N$$

at least one
machine out of
N fails

$$N=10: P \approx .01$$

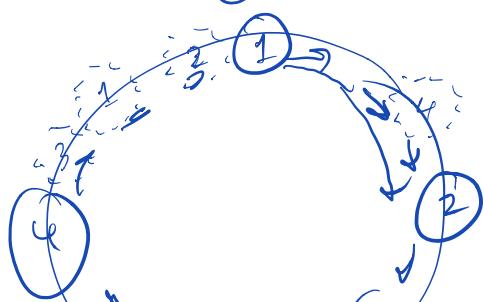
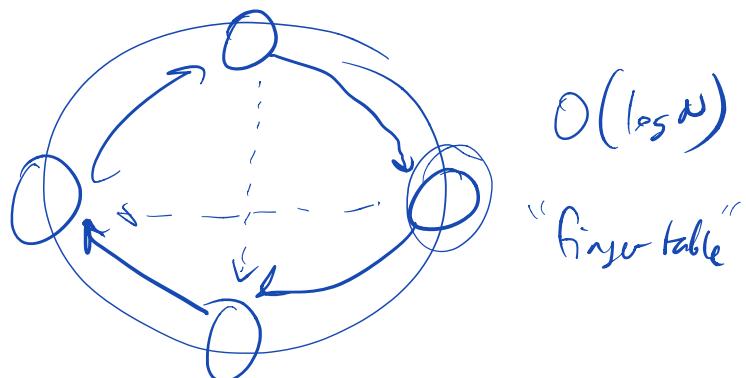
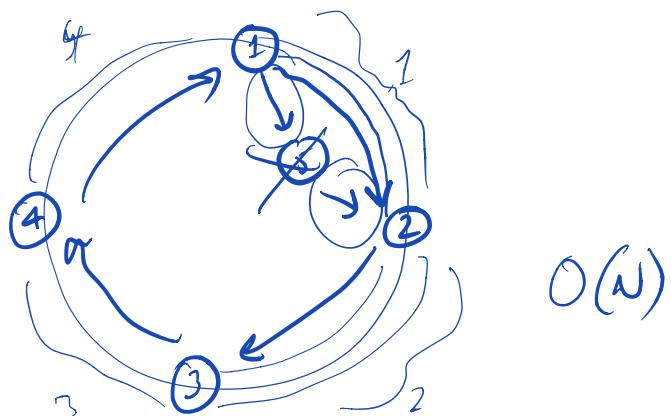
$$N=100: P \approx .1$$

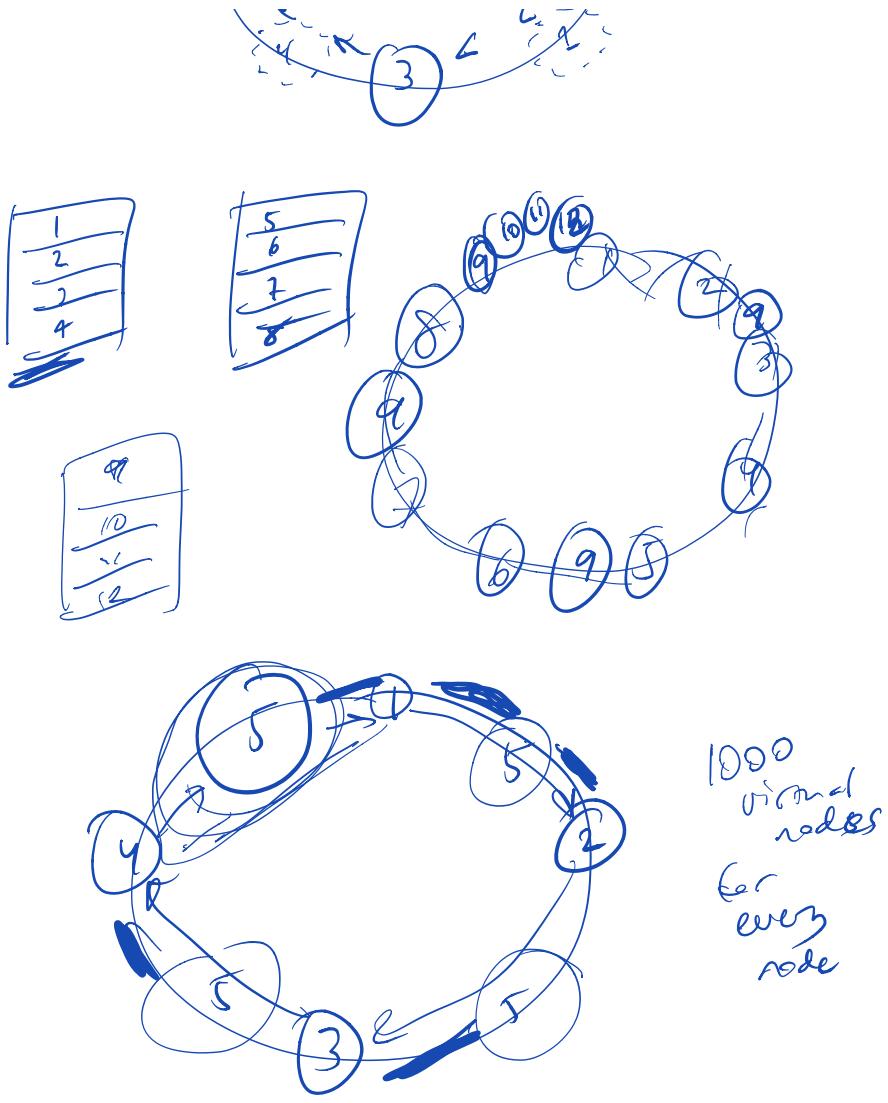
$$N=1000: P = 1 - (1-P)^{1000}$$
$$1 - .799^{1000}$$

$$\approx .63$$

Consistent hashing

Chord





Chord
consistent hashing