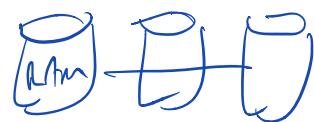
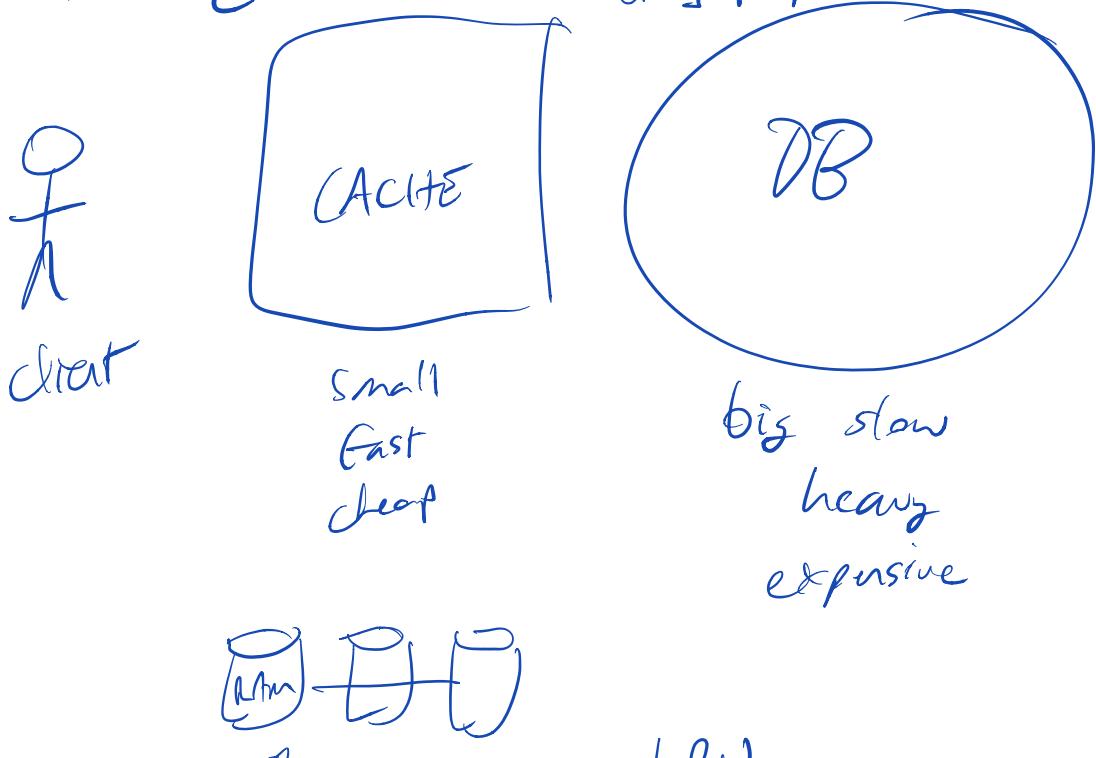


xavim@ac.upc.edu ← new  
unreg. people - send him a note

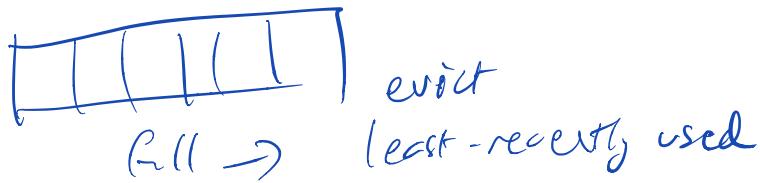


100%

LRU

replacement  
policies

past behavior  
is a good predictor  
of future behavior



A B C A B C A B C  
— — — 1 2 3 4 5 6 7 8 9

P	A	B	C
	6	7	5

A B C D E A B C D E  
 \_\_\_\_\_  
 | | | | |  
 6 7 8 9

MRU

LFU

Least Frequency Used ] Frequency

OPT

LRU  
 :  
 recency

"optimal"

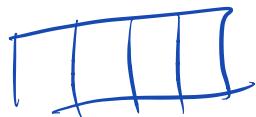
worst-case  
 size  
 of  
 cache

$k \times OPT$

best you can do

A B C D E A B C D E A B C D E

RAND



A B C A B C ...

X	3	4
	6	5

A B C D E A B C D E A B C D E

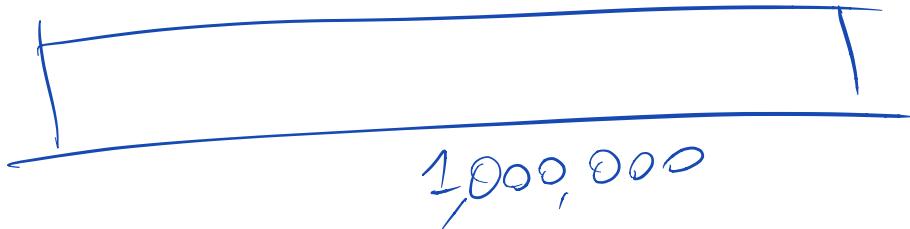
1 1 1 1 Y 1 1 1

$\frac{1}{4} \frac{1}{4} \frac{1}{4} \frac{1}{4} \dots \left(\frac{1}{4}\right)^n$

A	D	D	E
	6	7	5

A A C D A E B C

AS CD E 8888 F 8888 G 8888

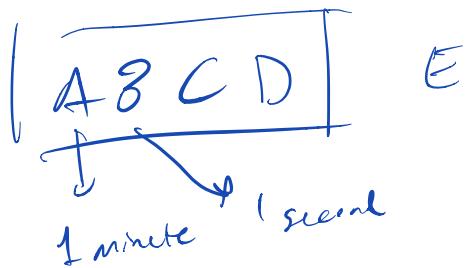
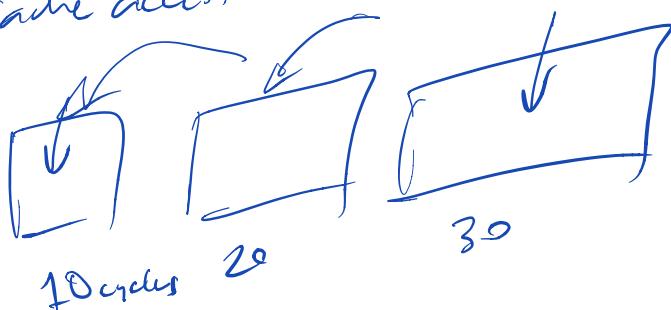


big cache = RAND ~ LRU  
in good case

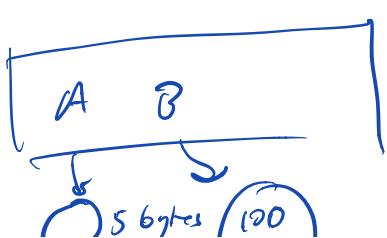
while avoiding  
pathological case

NUCA

non-uniform  
Cache access



assumes  
uniform  
cost

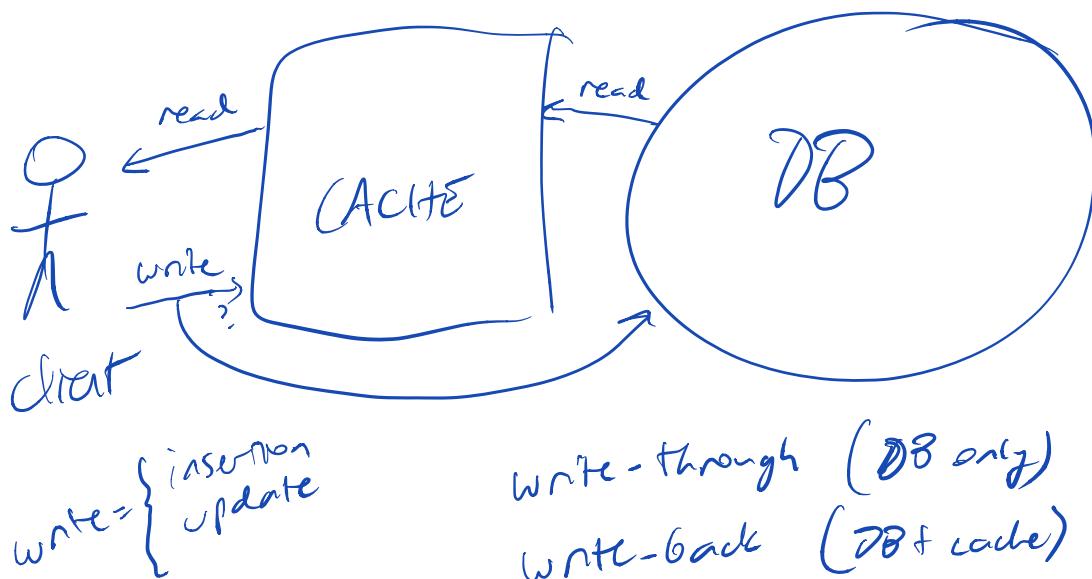
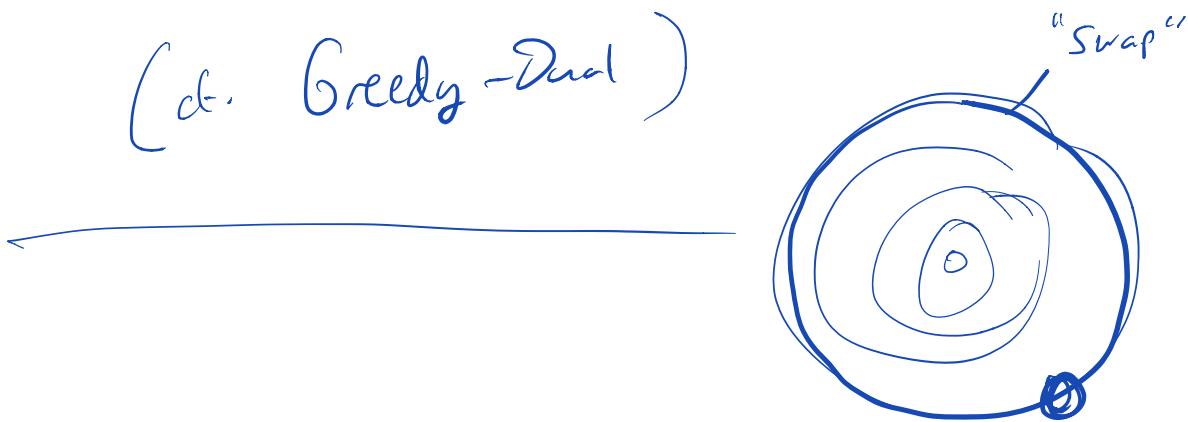


• uniform  
size

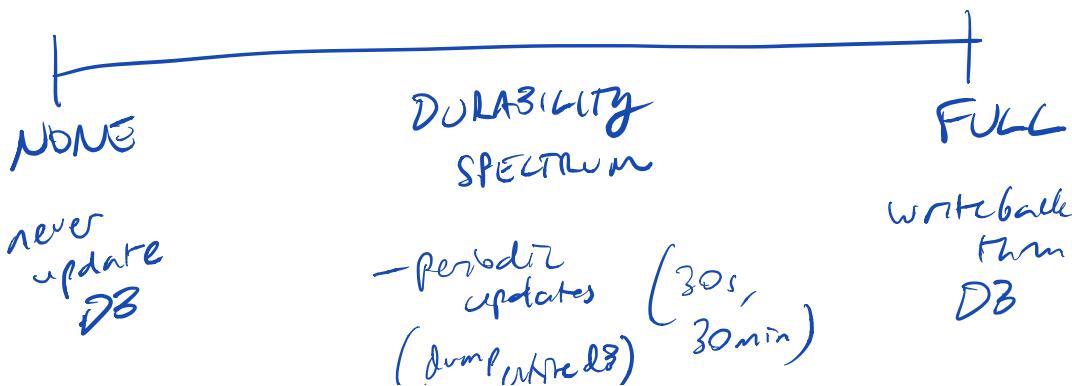
(MB)

Resource-aware replacement

(cf. Greedy-Dual)

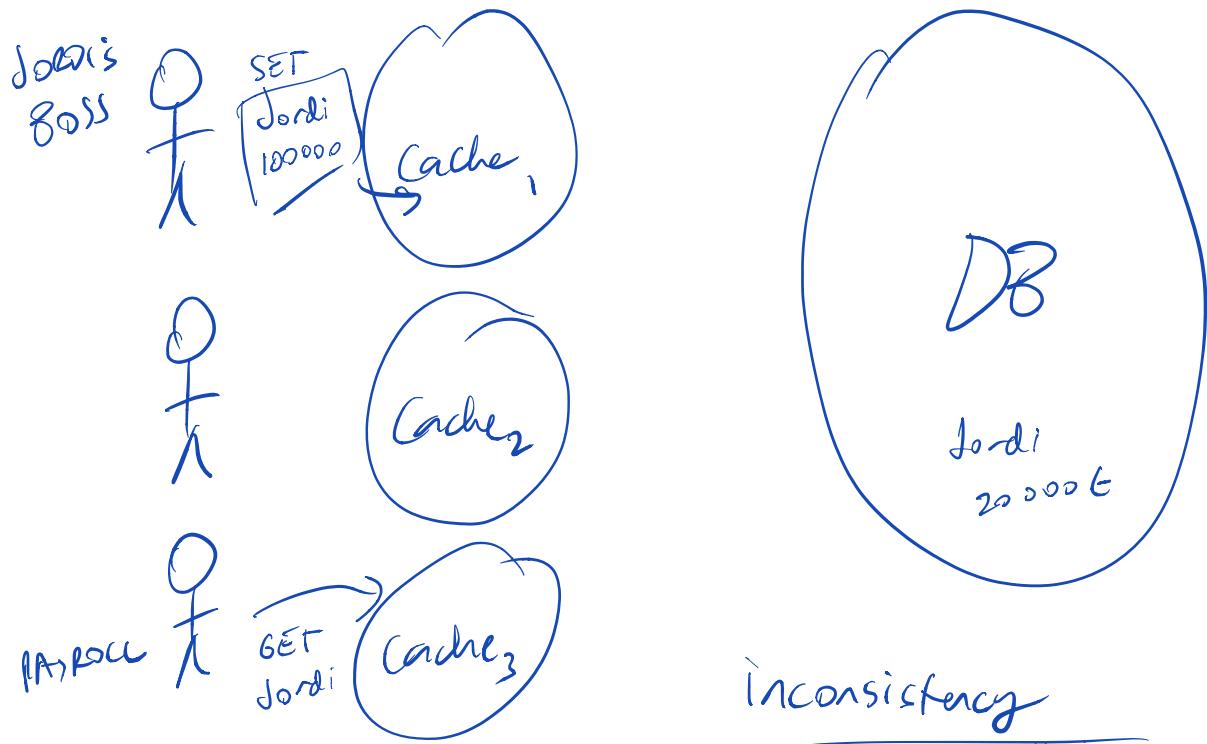
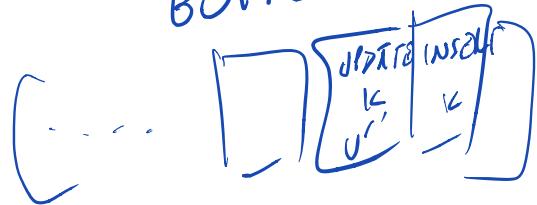


write-through (DB only)  
write-back (DB + cache)



- max #  
updates

## BUFFER



## Cache coherence

### "MESI"

modified

exclusive (only one copy)

shared (# copies > 1)

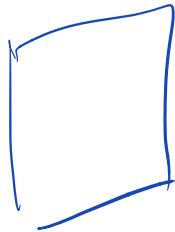
invalid

"lease" (30s)

SALARY
Jordi
100 000€

how long  
something  
can be  
kept in cache.

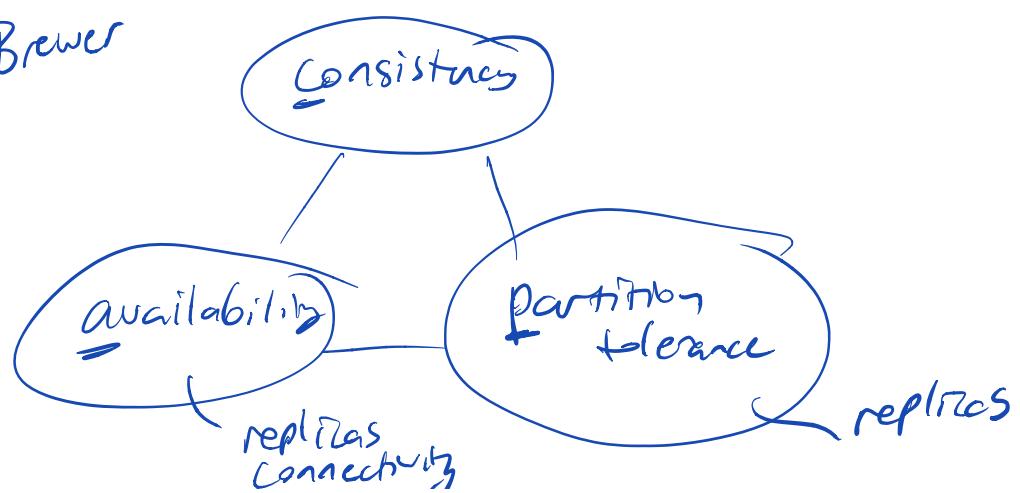
GET(SALARY,"Jordi", 1 week)



"eventual  
consistency"

"CAP theorem" — part two

Eric Brewer



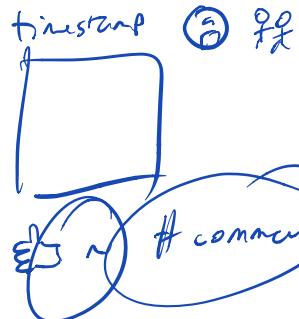
## Amazon

State  
views  
of  
# in stock

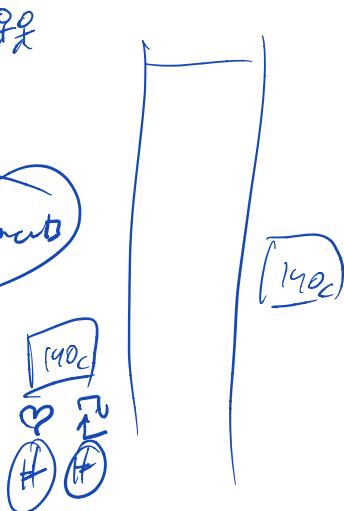
## eBay

state  
# of items/  
sold/not sold

## Facebook



## Twitter

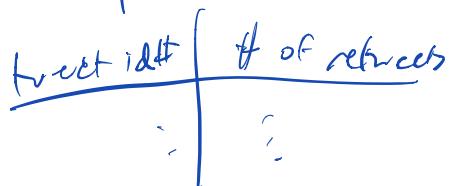
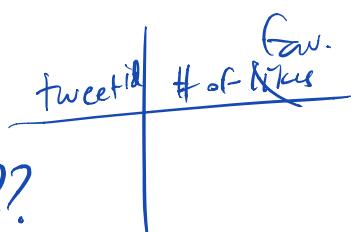


## NoSQL

- cost (DBs \$\$\$)
- eventual consistency OK
- flexible schema  
NO Schema

Schema rewrites??

### schema



— low latency (disk - o<sup>-t</sup>)

simple queries  
low latency needs ( $\sim 3$  orders of magn. faster)

— high scalability

Redis API

in-memory  
data structure  
store

dicts (hash tables)

strings

lists

sets

sorted set — range queries

bitmaps

5 [0001111000]

6 [110001111]

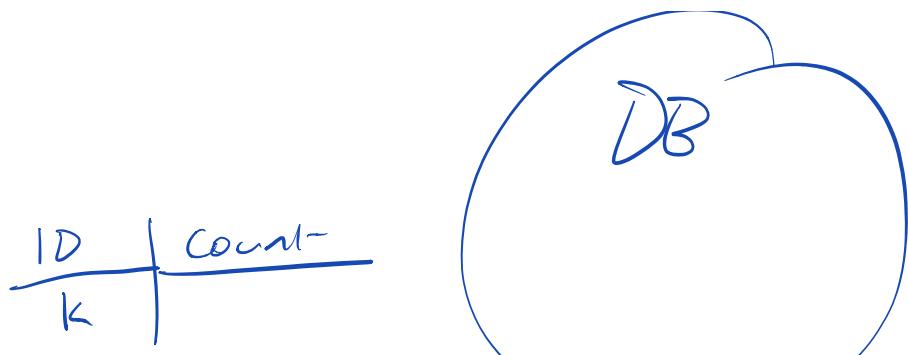
bitwise AND OR XOR...  
(1000001000)

—

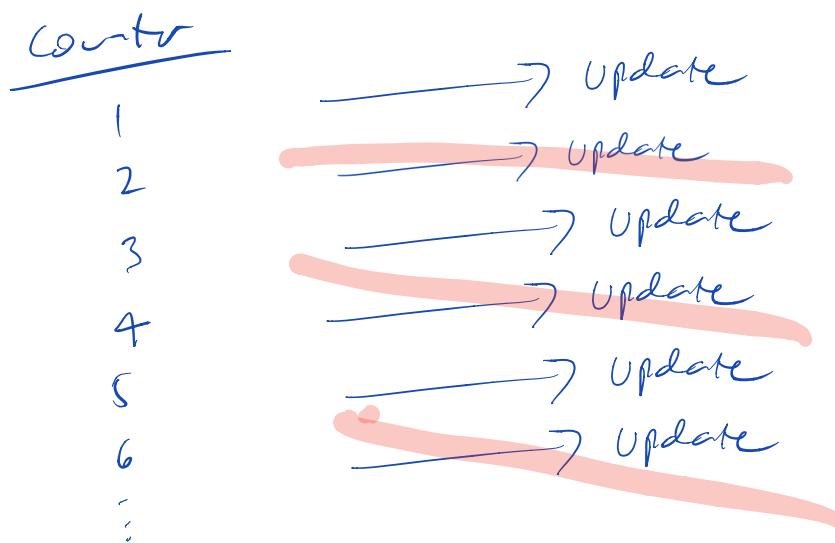
HyperLogLog

geospatial index (radius queries)

—

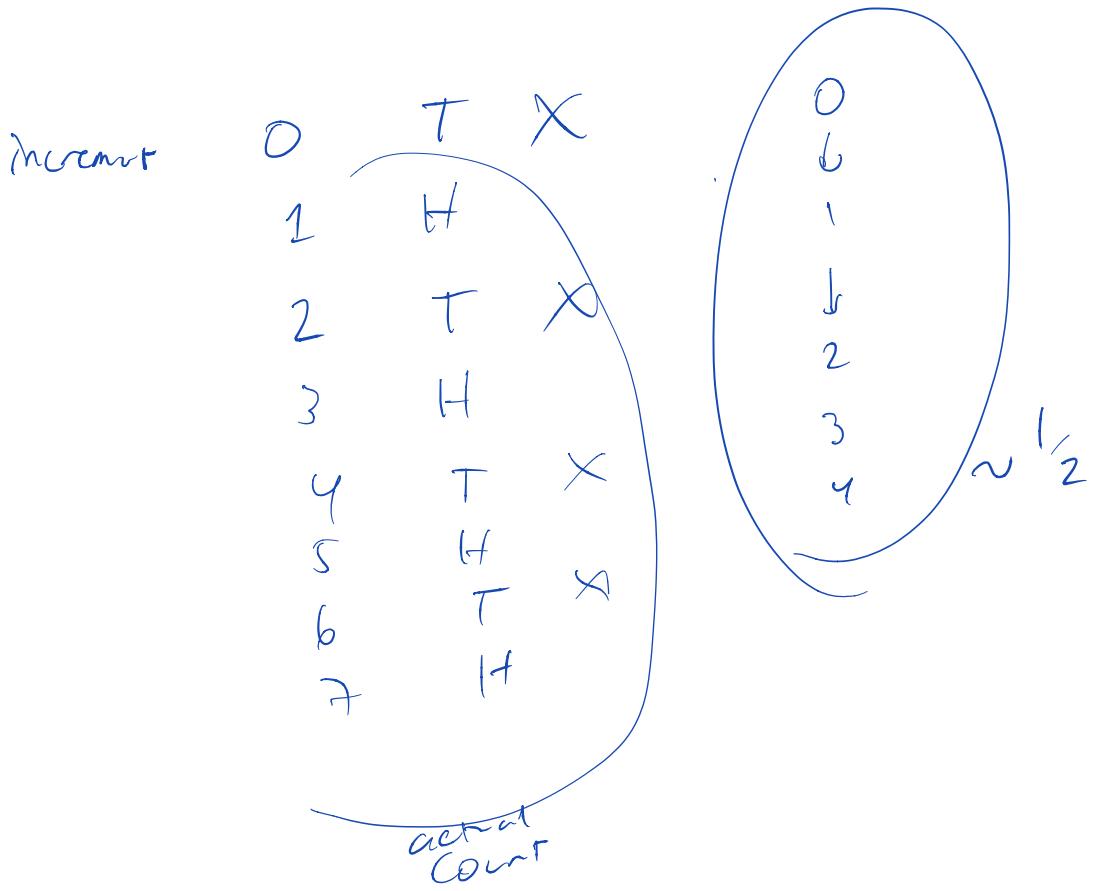


$c = \text{get}(k)$   
 $\text{set}(k, c+1)$

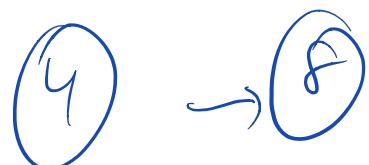


Robert Morris (97)

(dad of "inventor" of  
 Morris worm ~1988)



1	
2	$\frac{1}{2}$
4	$\frac{1}{4}$



$$\text{INC } \text{INC } \text{INC } \text{INC} \quad P(\text{success}) = \frac{1}{4}$$

$\theta \rightarrow \theta$

Q Q Q Q  
F F F F

Q Q Q Q  
F F F F

131072

262144

i

$\frac{1}{\sqrt{2}}$

2  $\frac{1}{2}$

F

4  $\frac{1}{4}$   
⋮

"approximate counting"

probabilistic data structures

trade accuracy for performance

$\text{add}(\text{item}, \text{set})$   
 $\text{check}(\text{item}, \text{set})$

set membership

$N$   
 $[00000000000000]$

$0 = \text{not in set}$   
 $1 = \text{in set}$

$\text{add}(k, \text{set})$   
 $[00001000000000]$

$\text{check}(k, \text{set})$

$0$  if  $1$ , return TRUE  
 $\text{else}$  FALSE

$\text{isEnd}(\text{env}, x)$   
 $[00001010101010]$

$N$

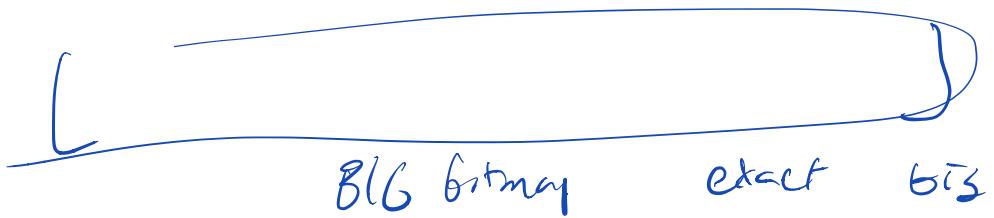
28 user

RLE  
run-length encoding

0	1000
1	3
0	1000

"TIFF"

# Bloom Filter



"FB Friend"

little bitmap X exact small ✓

NOT in SET — exact

in SET — ? maybe

0

No Friends

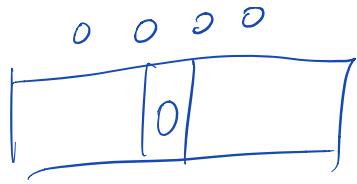
1

1 Friend

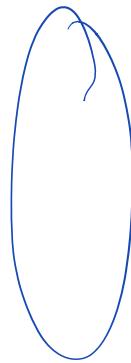
string	hash	?
"Serdil"	[1010]	?
"Paula"	[0011]	?
"Alp"	[0100]	X
"Alex"	[1011]	?

AND	0 1 0 0	0011
	-	-
	0 0 0 0	0011
	OR	-
	1 0 1 0	0011
	OR	-
	0 0 1 1	0011

Claudia



0 . . 1 - - -



# of Gatos

Cardinality

# of distinct things